

Reinforcement Learning

What is RL?

- 1) A problem
- 2) A community working on 1)
- 3) Methods produced by 2) which can be applicable to other problems

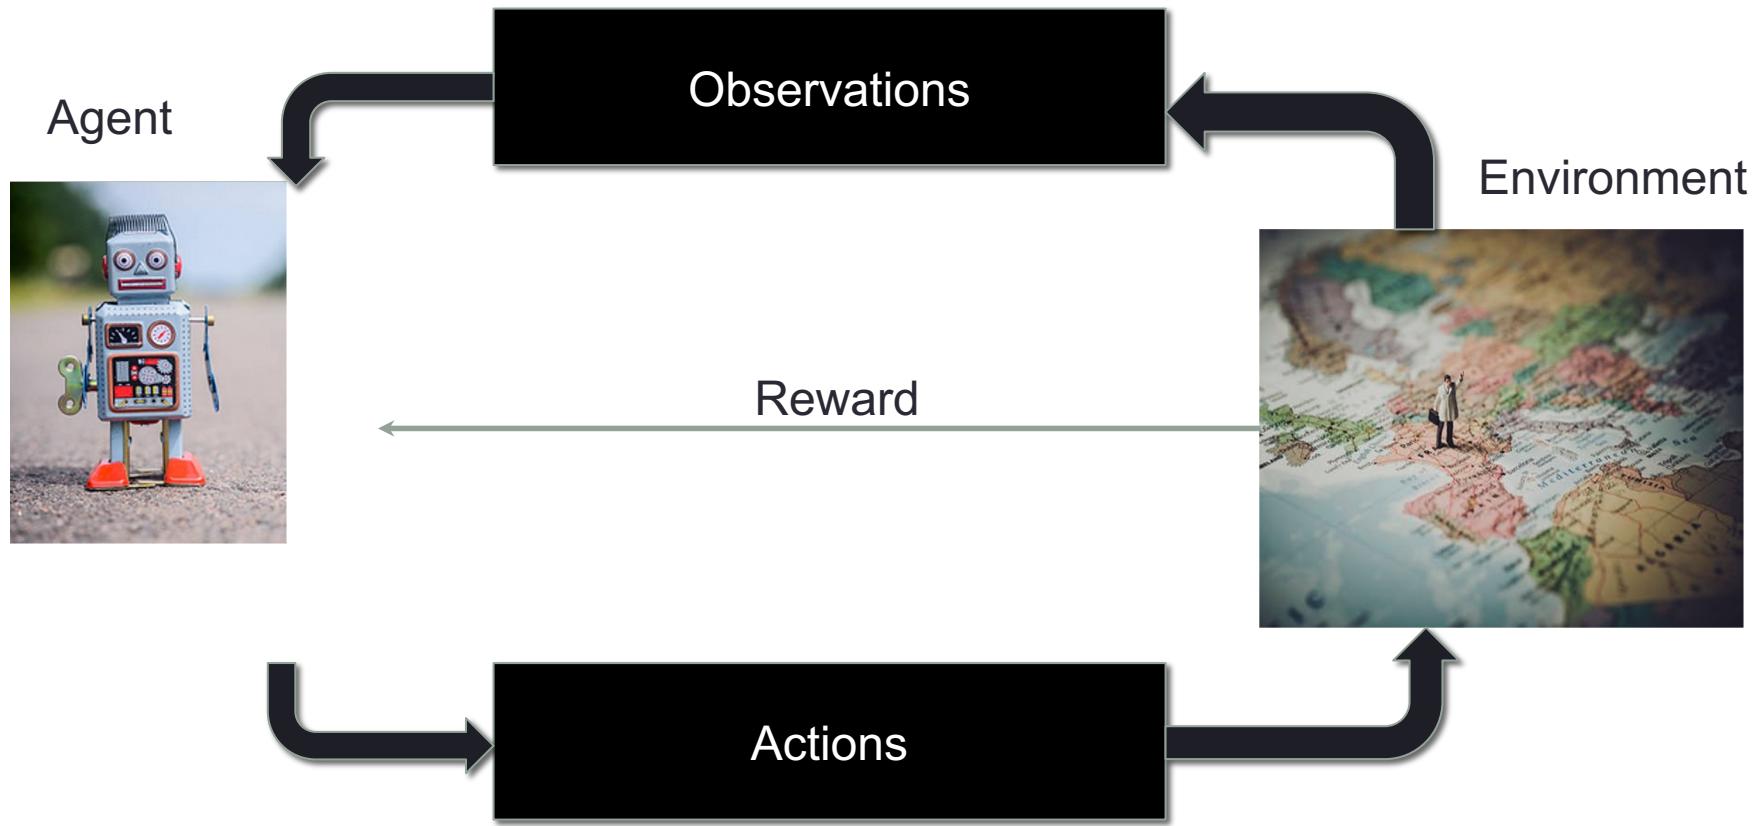


[Benjamin Van Roy](#)

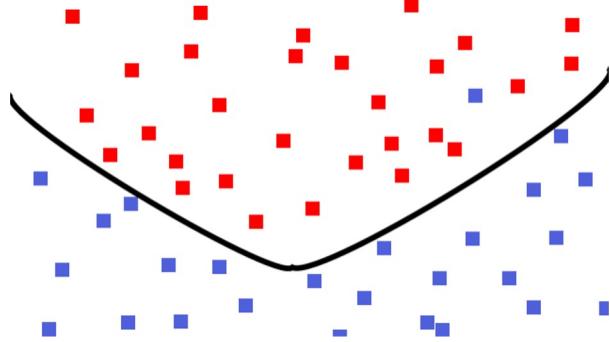
Professor at Stanford University; Research Lead at [DeepMind](#), Mountain View

Topic: Reinforcement Learning

RL problem



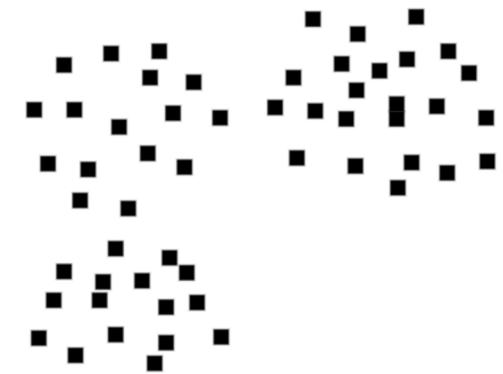
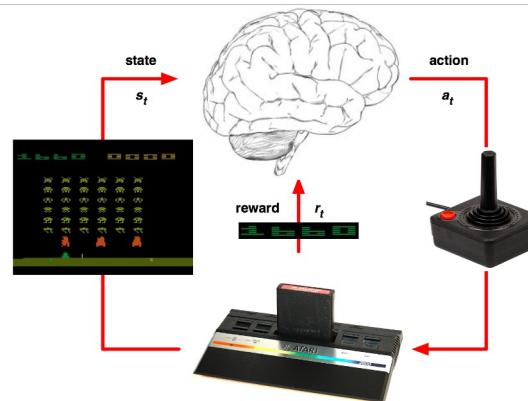
3 Modes of Learning



Supervised Learning



Reinforcement Learning

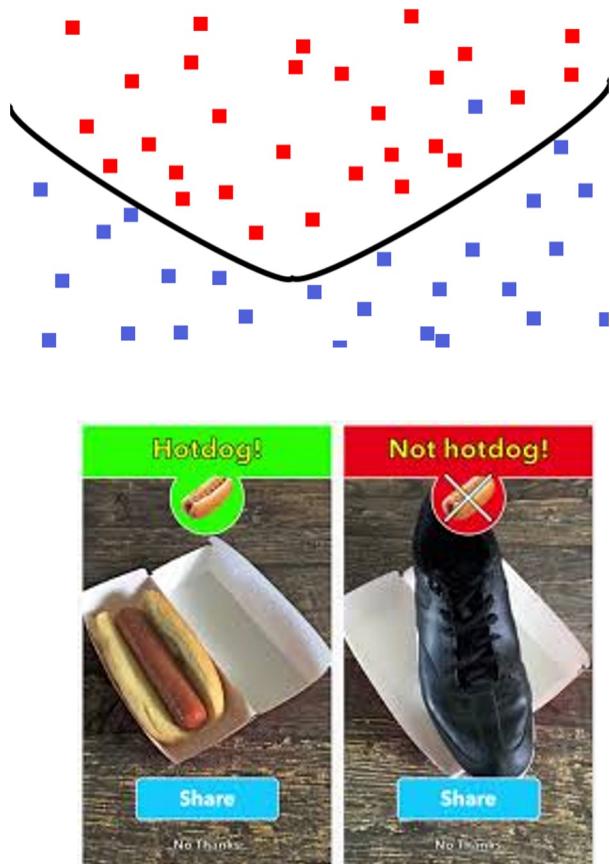


Unsupervised Learning



3 Modes of Learning

Supervised Learning

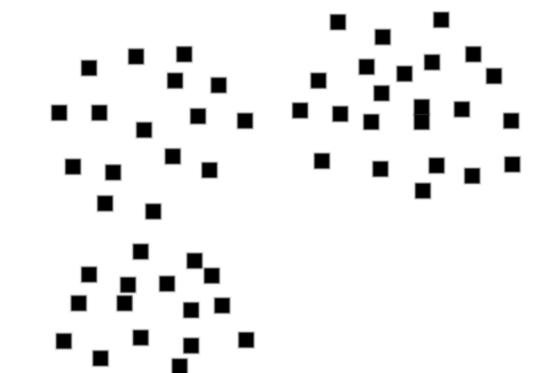


- Observe:
 - $(x_1, y_1), (x_2, y_2), \dots$
- Objective:
 - Input an unseen x_{new}
 - What is y_{new} ?

3 Modes of Learning

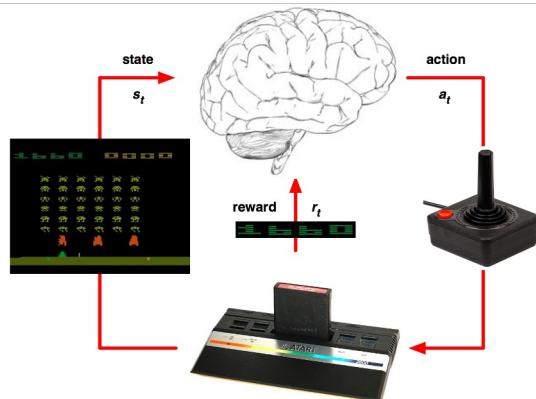
Unsupervised Learning

- Observe:
 - $x_1, x_2, x_3, x_4, \dots$
- Objective:
 - What is $P(x)$?
 - What is a *good* representation of x ?
 - What can we learn from $P(x)$?



3 Modes of Learning

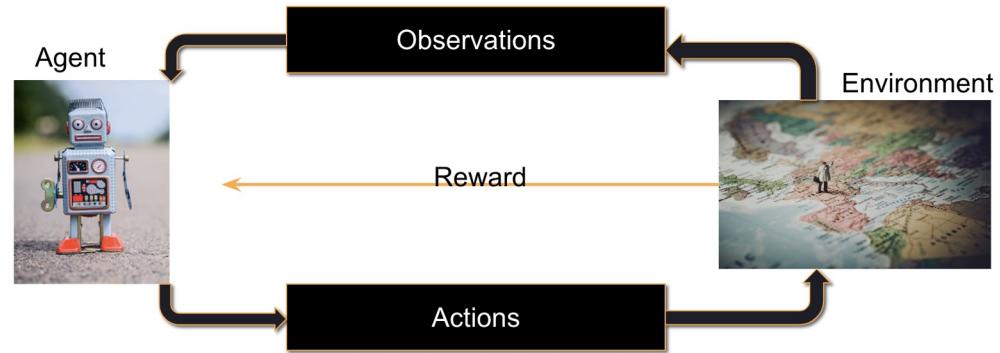
Reinforcement Learning (RL)



- Observe:
 - The states (x_1, x_2, x_3, \dots)
 - The reward (r_1, r_2, r_3, \dots)
- Can also take actions
 - a_1, a_2, a_3, \dots
- What are the best actions?
 - Such that we will receive highest accumulative rewards

Difference between RL and other modes of learning

- Sequential decisions
- You have a goal vs You have means to get there
- No concept of “training set” and “test set”
- “Passive” vs “Active” learning



RL and Artificial General Intelligence



Yann LeCun

March 14, 2016 · 0

Follow

Statement from a Slashdot post about the AlphaGo victory: "We know now that we don't need any big new breakthroughs to get to true AI"

That is completely, utterly, ridiculously wrong.

As I've said in previous statements: most of human and animal learning is unsupervised learning. If intelligence was a cake, unsupervised learning would be the cake, supervised learning would be the icing on the cake, and reinforcement learning would be the cherry on the cake. We know how to make the icing and the cherry, but we don't know how to make the cake.

We need to solve the unsupervised learning problem before we can even think of getting to true AI. And that's just an obstacle we know about. What about all the ones we don't know about?

#deeplearning #AI #AlphaGo

This is before he coined the term self-supervised learning for supervised learning

The Yann Lecunn's cake

Y. LeCun

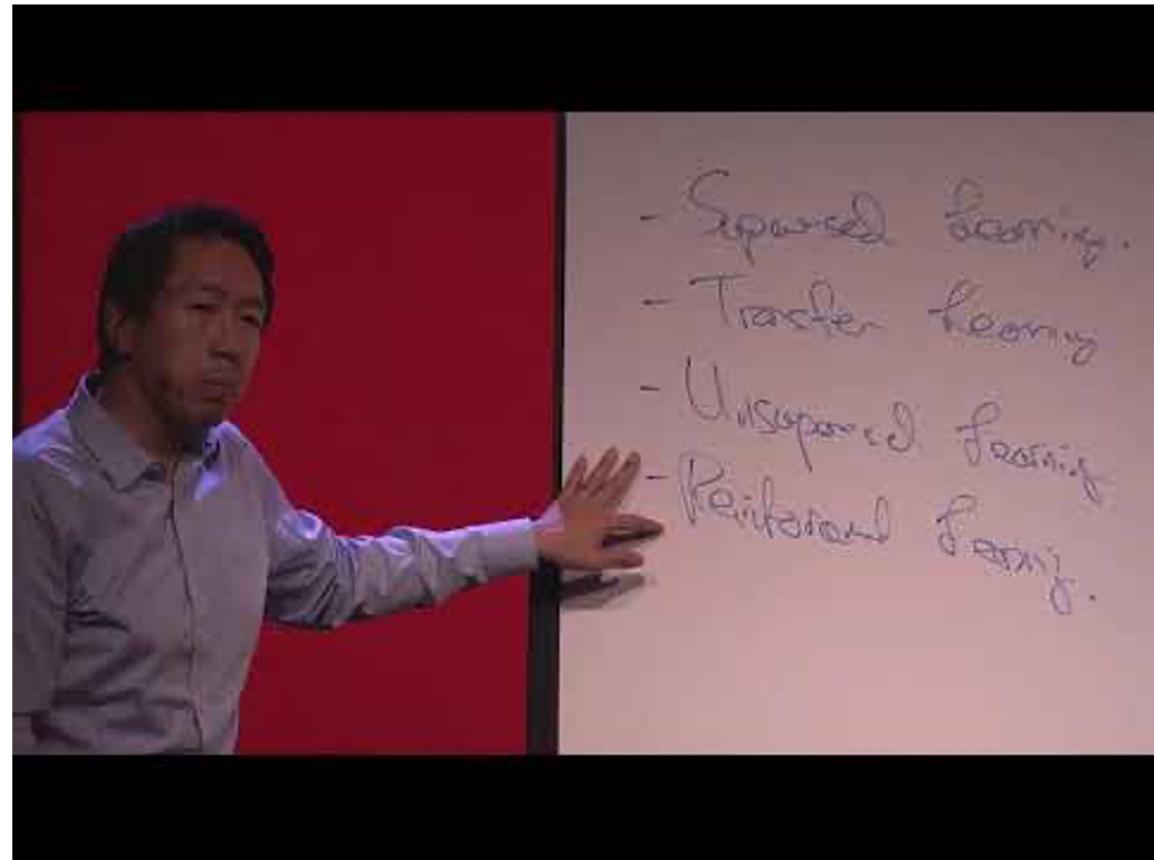
How Much Information is the Machine Given during Learning?

- ▶ “Pure” Reinforcement Learning (**cherry**)
 - ▶ The machine predicts a scalar reward given once in a while.
- ▶ **A few bits for some samples**
- ▶ Supervised Learning (**icing**)
 - ▶ The machine predicts a category or a few numbers for each input
 - ▶ Predicting human-supplied data
 - ▶ **10→10,000 bits per sample**
- ▶ Self-Supervised Learning (**cake génoise**)
 - ▶ The machine predicts any part of its input for any observed part.
 - ▶ Predicts future frames in videos
 - ▶ **Millions of bits per sample**



RL and \$\$\$

“The excitement and PR hype behind reinforcement learning is a bit disproportionate relative to the economic value it's creating today” - Andrew Ng



RL use cases

Go, chess, starcraft, dota, poker

Finance

(<https://www.jpmorgan.com/global/LOXM>)...

ChatGPT

Robotics...but...



<https://research.googleblog.com/2016/03/deep-learning-for-robots-learning-from.html>

<https://towardsdatascience.com/applications-of-reinforcement-learning-in-real-world-1a94955bcd12>

<https://www.oreilly.com/ideas/practical-applications-of-reinforcement-learning-in-industry>

RL use cases

Data center and resource management (<https://people.csail.mit.edu/alizadeh/papers/deeprm-hotnets16.pdf>)
System configuration <http://ranger.uta.edu/~jrao/papers/ICDCS09.pdf> DRAM controller
<https://ieeexplore.ieee.org/abstract/document/4556714/> Recommender (Bandits)
<https://people.cs.umass.edu/~pthomas/papers/Barto2017.pdf>)

Ads bidding (<https://arxiv.org/abs/1701.02490>)

Chemistry (<https://pubs.acs.org/doi/full/10.1021/acscentsci.7b00492>)

Some other tasks that use algorithms from RL to help perform model training (autoML, REINFORCE)

[RETURN TO ISSUE](#) | < PREV **ARTICLE** NEXT >

Optimizing Chemical Reactions with Deep Reinforcement Learning

Zhenpeng Zhou[†] , Xiaocheng Li[‡] and Richard N. Zare^{*†} 

[View Author Information](#) ▾

 [Cite This: ACS Cent. Sci.](#) 2017, 3, 12, 1337-1344

Publication Date: December 15, 2017 

<https://doi.org/10.1021/acscentsci.7b00492>

Copyright © 2017 American Chemical Society

[RIGHTS & PERMISSIONS](#)  [ACS AuthorChoice](#)

Article Views
13256

Altmetric
20

Citations
9

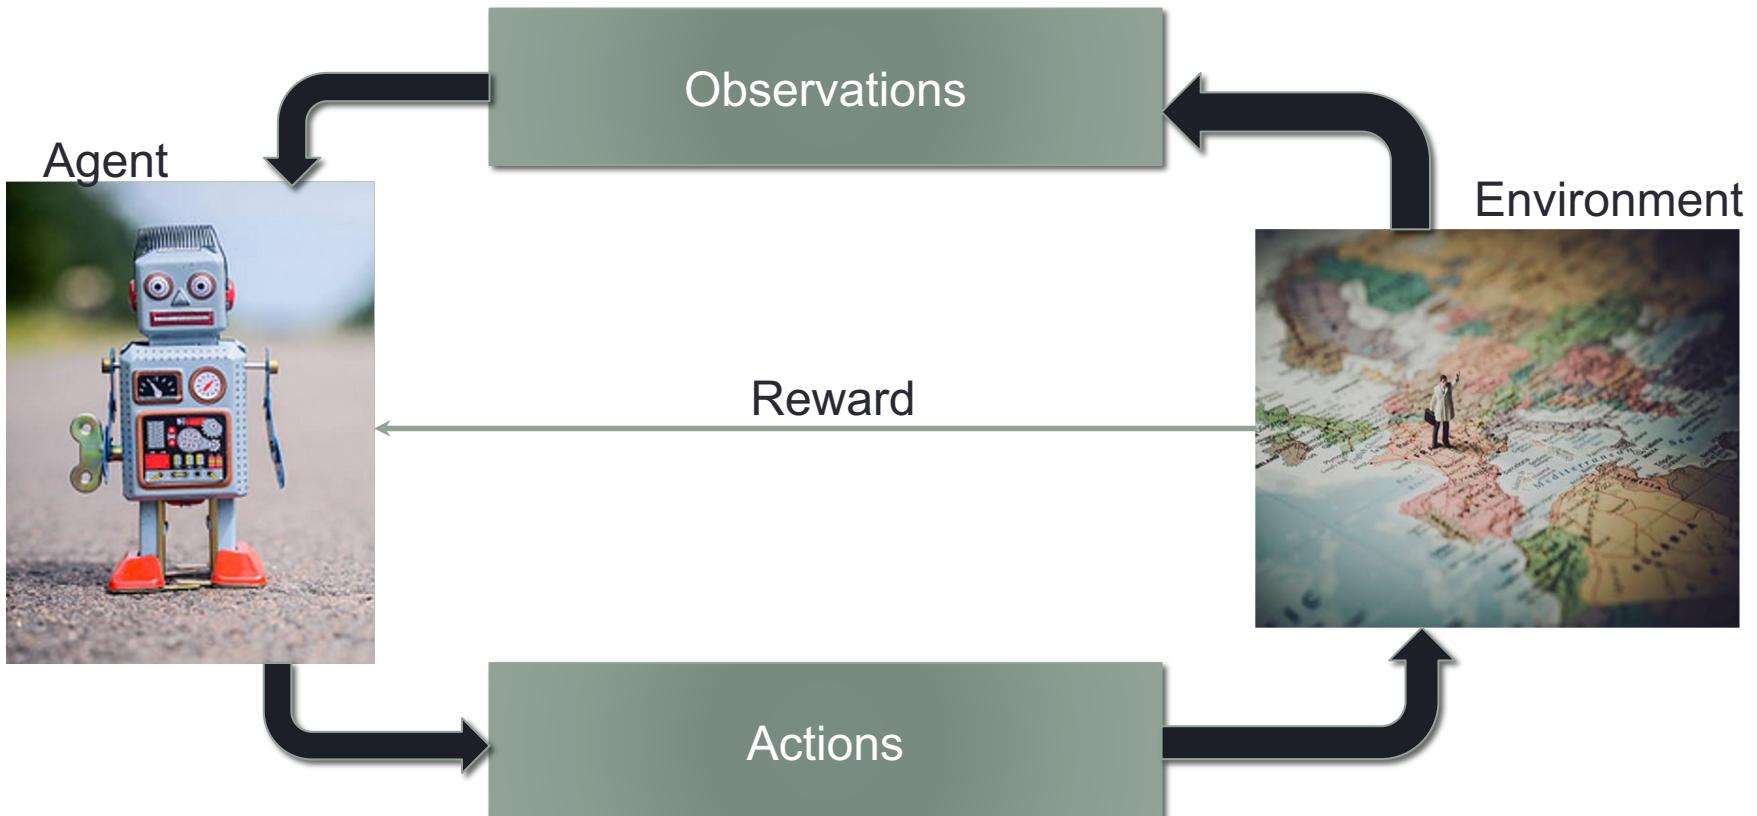
Share  Add to  Export 
[LEARN ABOUT THESE METRICS](#)



ACS Central Science

 [PDF \(3 MB\)](#)

RL framework



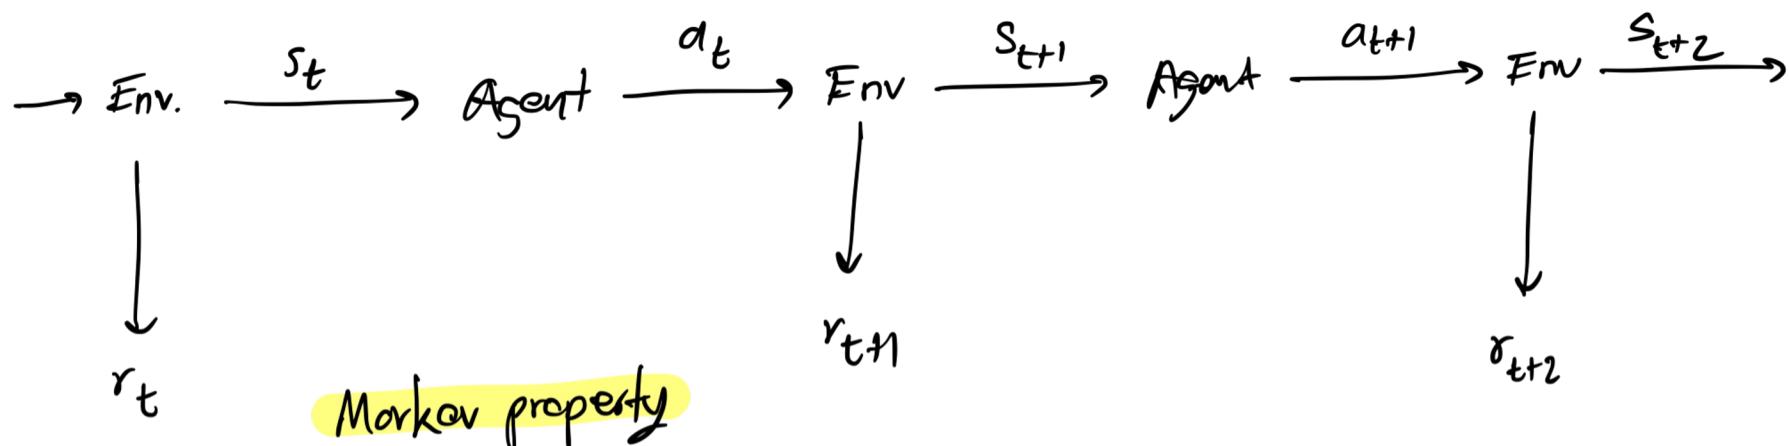
Learning through trial and error

RL framework

Reward (r_t)

State (s_t)

Action (a_t)



$$a_t = A(s_t)$$

$$r_{t+1} = R(s_t, a_t)$$

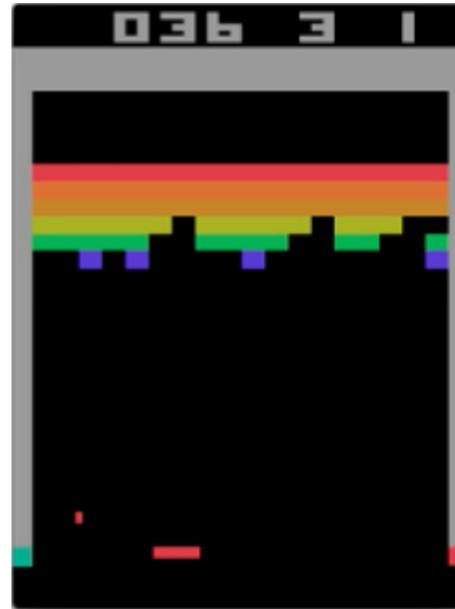
$$s_{t+1} = S(s_t, a_t)$$

Rewards-based learning

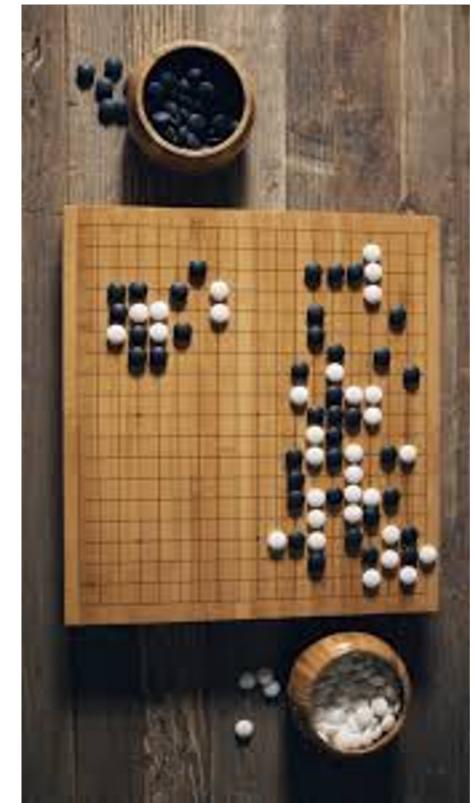
- Maximise the rewards
- Can we design any desired behaviour with reward?



$R_t = \Delta \text{distance}$



$R_t = \text{score}$



$$R_T = \begin{cases} 1, & \text{win} \\ -1, & \text{lose} \end{cases}$$

The Environment

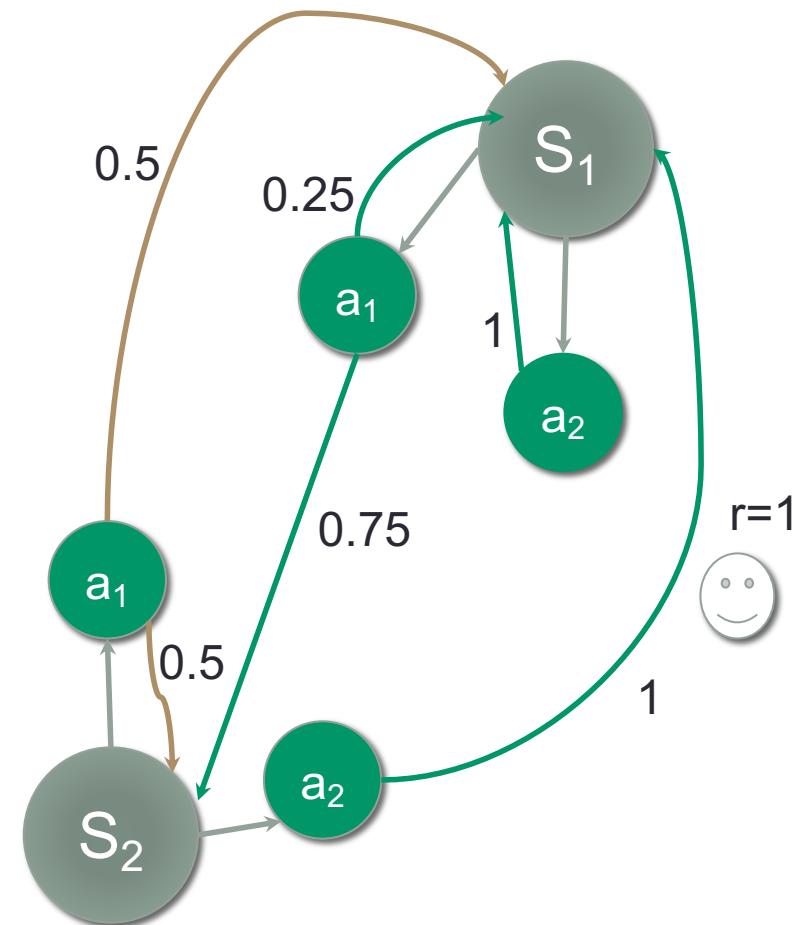
How can we model the environment?

Markov Decision Process (MDP)

- **S,A,P,R, γ**
- **S** – Set of states
- **A** – Set of actions
- **P** – Transition between states given an action

$$P_{s,s'}^a = \text{Prob}[s_{t+1} = s' \mid s_t = s, a_t = a]$$

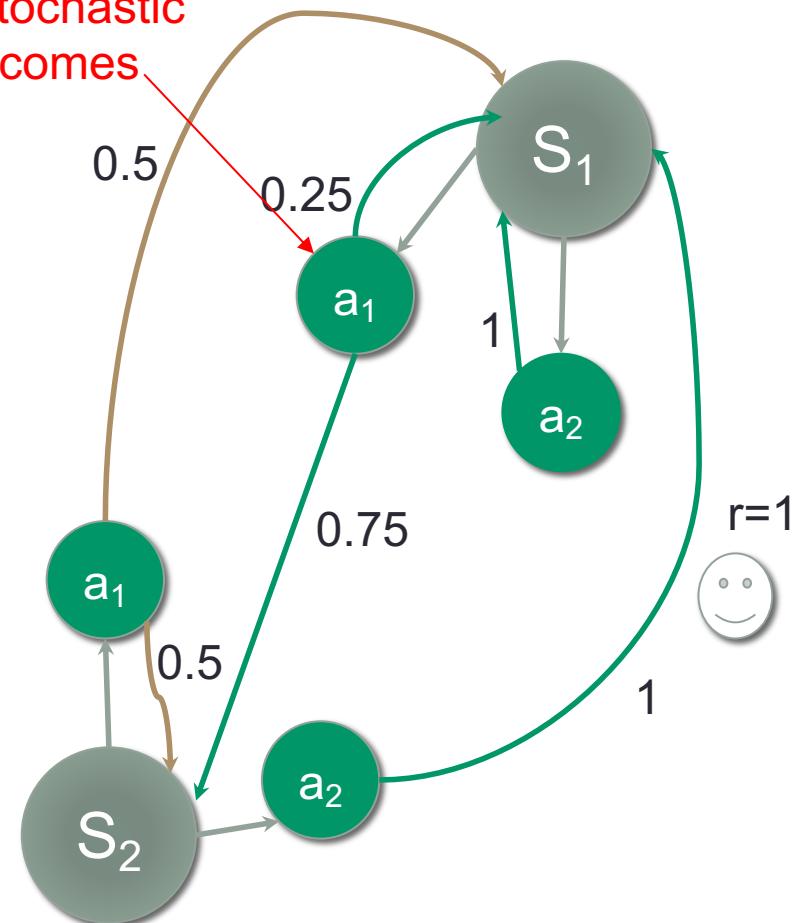
- **R** – Rewards associated with actions and states
- **γ** – Discount factor



Markov Decision Process (MDP)

- **S,A,P,R, γ**
- **S** – Set of states
- **A** – Set of actions
- **P** – Transition between states given an action
 $P_{s,s'}^a = \text{Prob}[s_{t+1} = s' | s_t = s, a_t = a]$
- **R** – Rewards associated with actions and states
- **γ** – Discount factor

Environment could be stochastic
One action, multiple outcomes



Markov Property

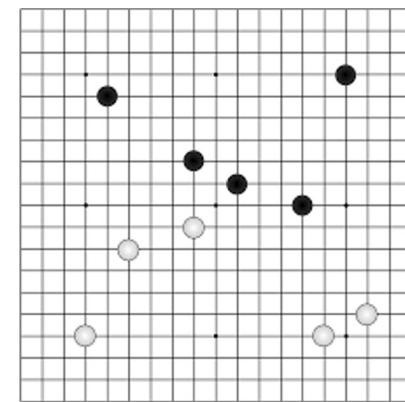
$$p(s_{t+1} | s_t, a_t)$$

- s_{t+1} depends only on s_t
- not s_{t-1} , not anything before
- this simplifies our situation!

But, is it true in every case?

- It depends on your observed state
 - Fully observable state 
 - Partially observable state 

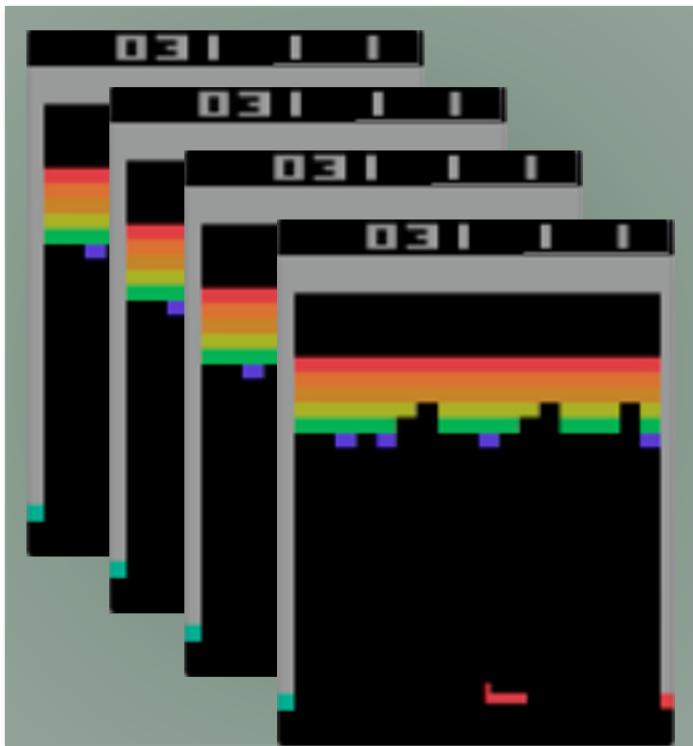
Fog of war



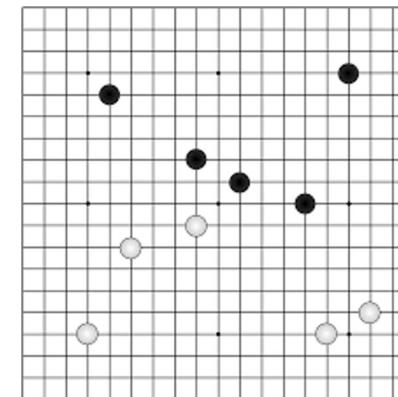
Fully observable

Fully Observable State

- Fully observable state: All information from the past is captured in the current state



For Go, a board position
For simple video games, stack multiple frames



The Agent

Policy

- Policy = a mapping from a state to an action
- Objective of RL is to find the “*optimal*” policy!

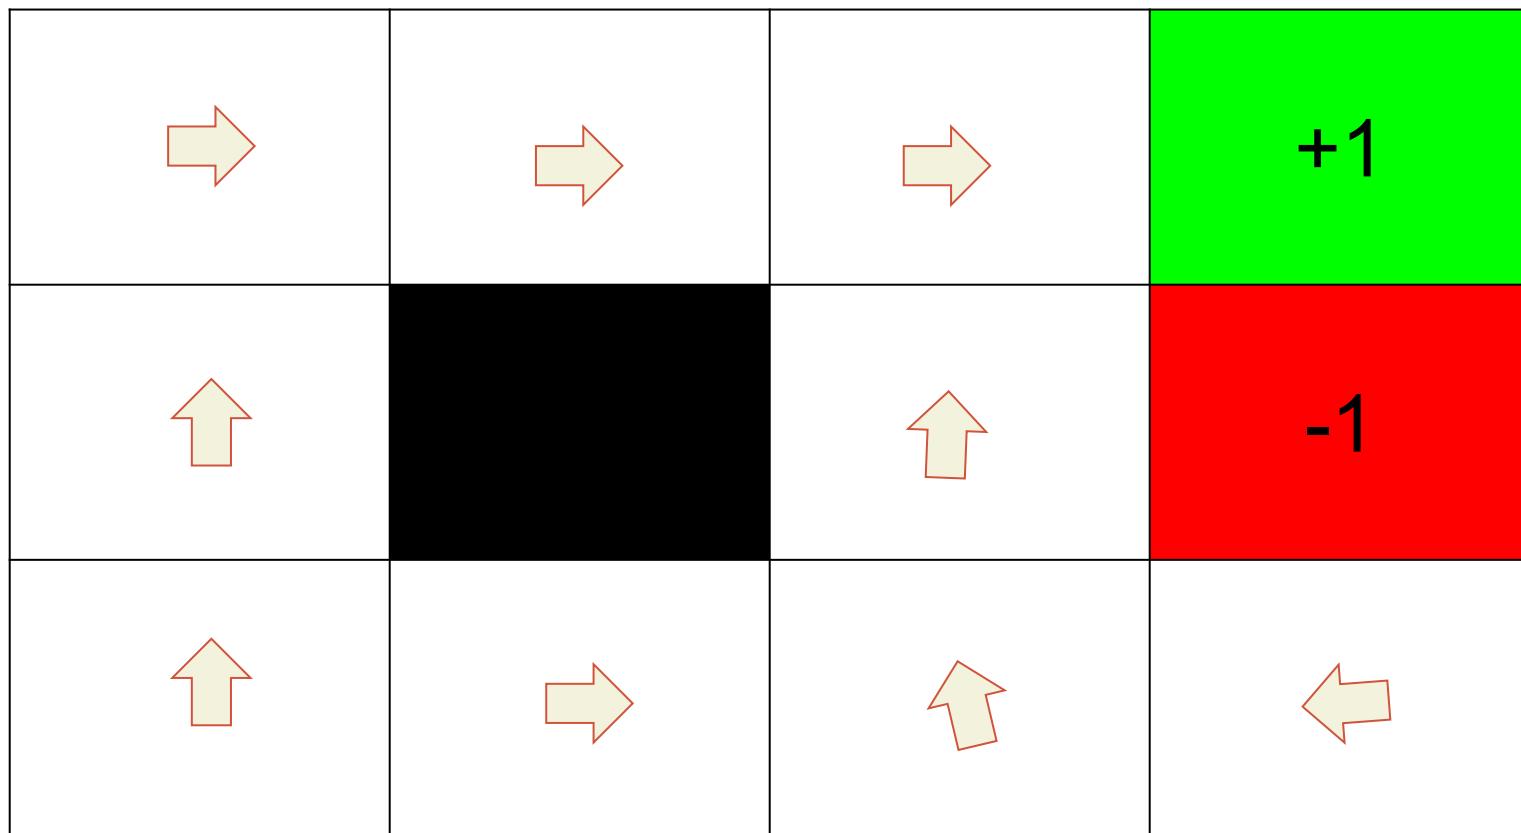
$$a = \pi(s)$$

Can be either
deterministic or
stochastic

$$a \sim \pi(s)$$

Policy

Example: tabular policy

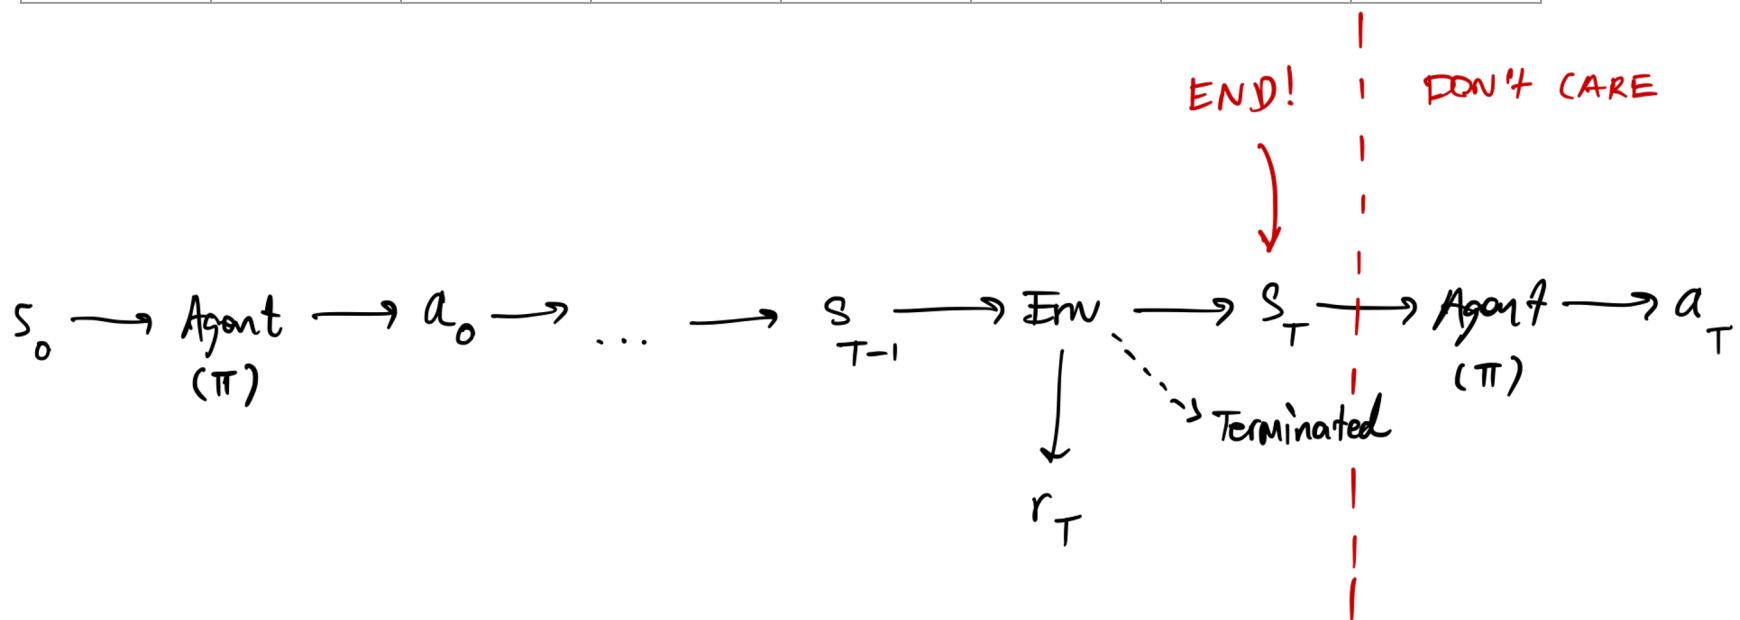


Learning

How do we find the best policy ?

Rollout (Our data)

Time	0	1	2	3	...	T-1	T	
S	s_0	s_1	s_2	s_3	...	s_{T-1}	s_T	Don't care
A	a_0	a_1	a_2	a_3	...	a_{T-1}	a_T	
R	r_0	r_1	r_2	r_3	...	r_{T-1}	r_T	
Done	0	0	0	0	...	0	1	



Return (Cumulative rewards)

- Return = cumulative rewards with discount

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots = \sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'}$$



$$r_t = 0$$



$$r_{t+10} = 1$$



$$r_{t+34} = -1$$

What is learning?

- Use data to find/search for the best policy

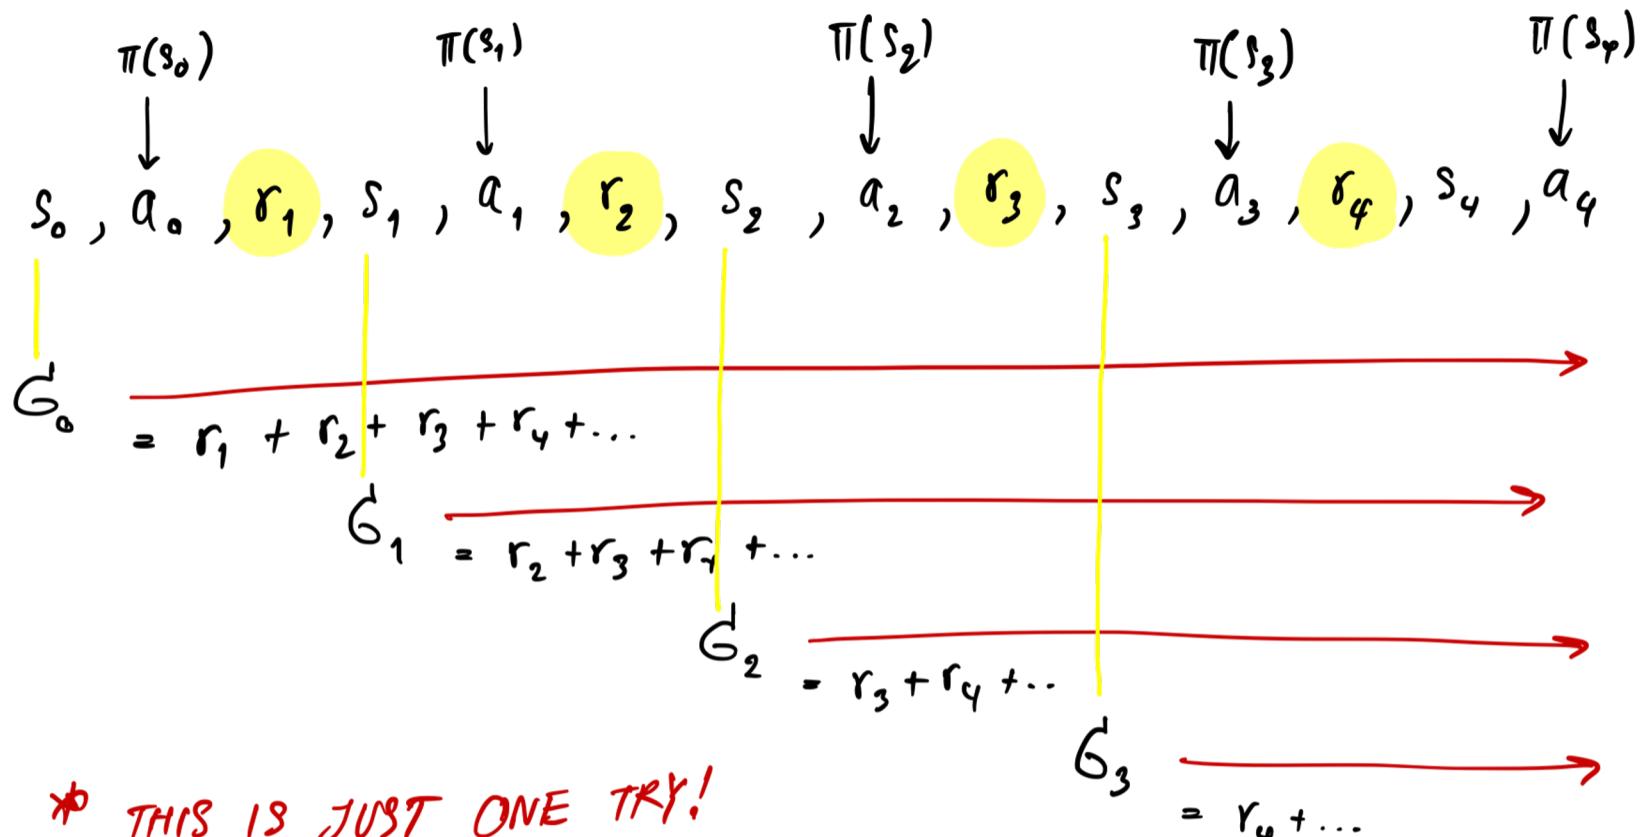
What is the best policy?

- Policy that give us the highest expected return!

$$G = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$$

$$J(\pi) = E_{\pi}[G]$$

Return under a policy (G^π)



* THIS IS JUST ONE TRY!
ANOTHER MIGHT YIELD DIFFERENTLY

Expected Return

How good “on average” is our return?

$$E_{\pi} [G_0]$$

statistical expectation

following policy π

For value G_0

- Start
- Play with π
- $G_0 = \sum_{t=0}^T r_t$
- Retry many times (∞)
- Average G_0

A naive learning method

1. Initialise a policy randomly
2. Evaluate the policy by running that policy multiple times
 - a. which we then collect the returns of all the runs
3. Randomly initialise another policy
4. Evaluate the new policy
5. Keep the policy that have a higher expected return
6. Repeat 3-5

Intuitive. But very inefficient!

How to make it more efficient?

Categories of RL methods

Value-based vs Policy-based

Model-based vs Model-free

On-policy vs Off-policy

Trade-off on a spectrum

Valued-based RL

Q-learning algorithm

- Let's define a state value as
 - Expectation of the return after visit s and follow π

$$V^\pi(s) = E_\pi[G_t \mid s_t = s]$$

- Let's define a state-action value (Q-value) as
 - Expectation of the return after visit a state s, take action a

$$Q^\pi(s, a) = E_\pi[G_t \mid s_t = s, a_t = a]$$

$$V^\pi(s) = E_{\pi(s)}[Q^\pi(s, \pi(s))]$$

Q-learning algorithm

- There exist an optimal value function associate with an optimal policy,

$$V^*(s) = \max_{\pi} V^{\pi}(s) \quad \forall s \in S$$

- The optimal policy is the policy that achieves the highest value for every state

Q-learning algorithm

- It follows that

$$V^*(s) = \max_a [Q^*(s, a)]$$

- and ..

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- Optimal actions can be found indirectly through Q-value

Value-based learning

2 Steps

Improve Q/V

Improve the policy

How to learn Q/V?

Monte-Carlo

Bootstrap

How to improve the policy?

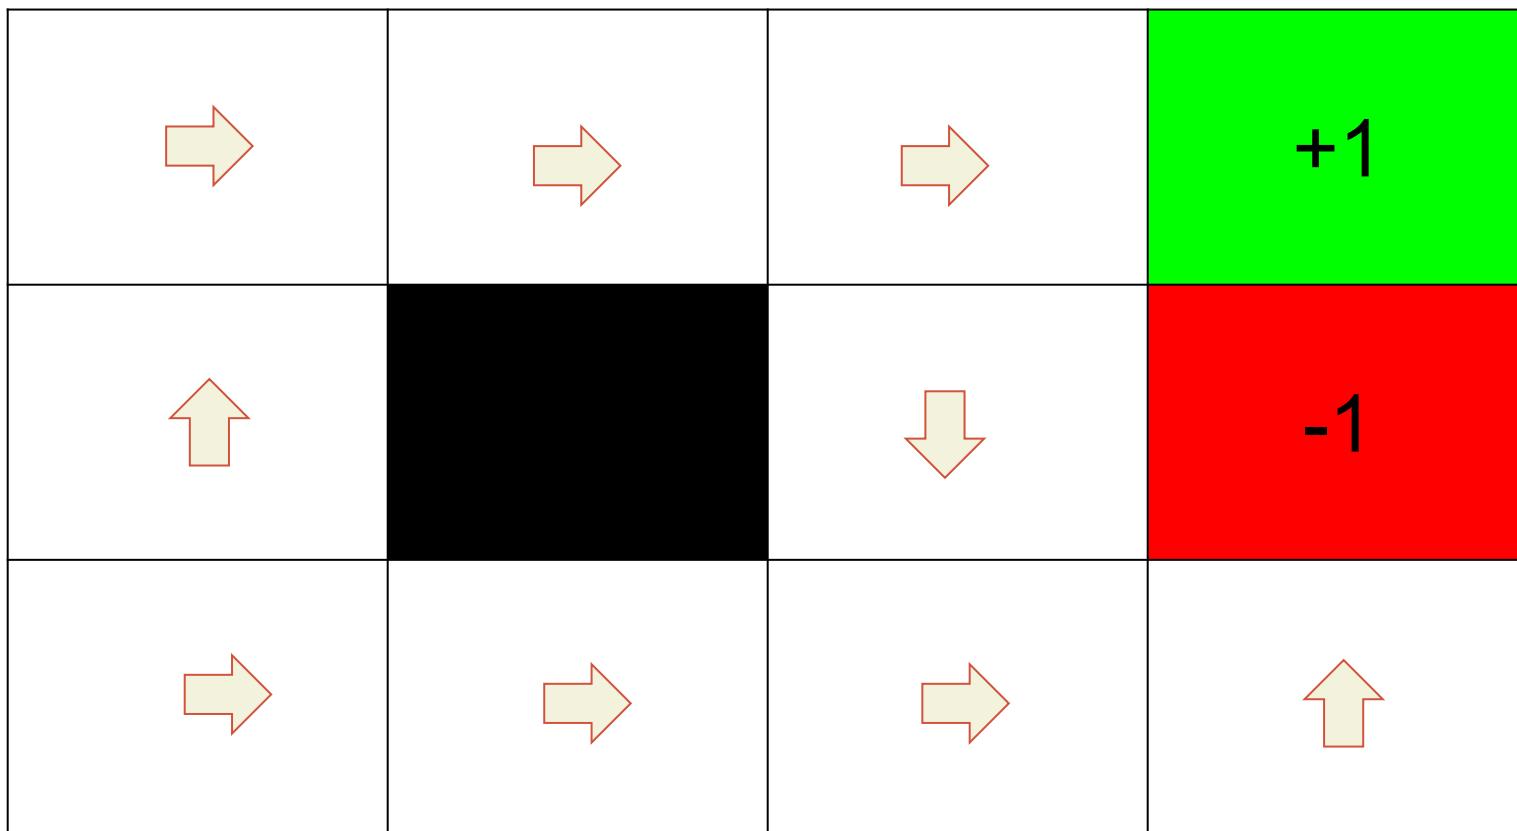
Follows the best Q/V

Example, Tabular Q-learning

			+1
			-1

Monte-Carlo Estimator

- Initialise π

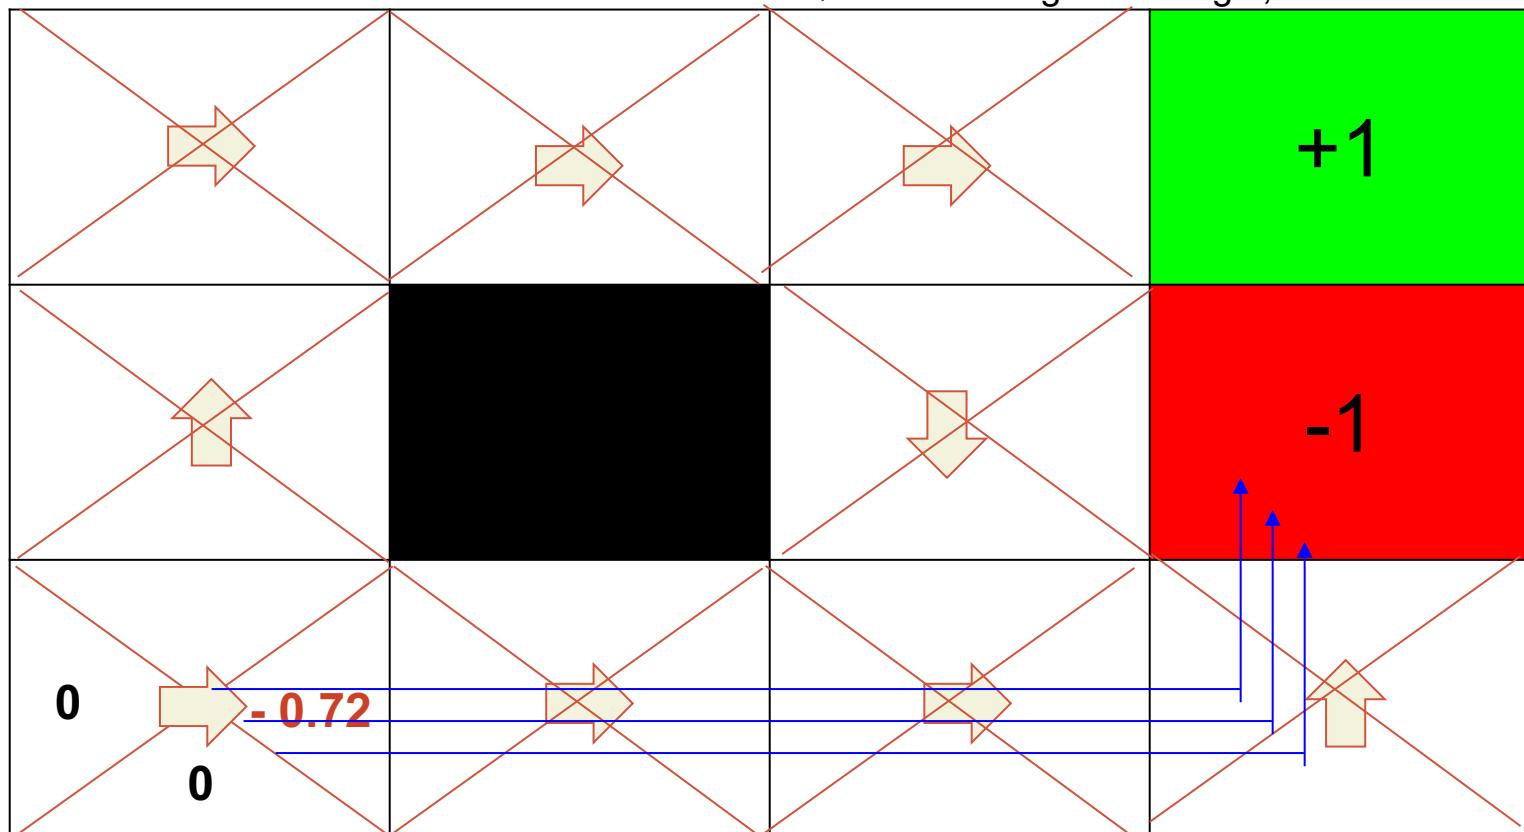


Monte-Carlo Estimator

$$Q^\pi(s, a) \approx r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^n r_n$$

$$\gamma = 0.9$$

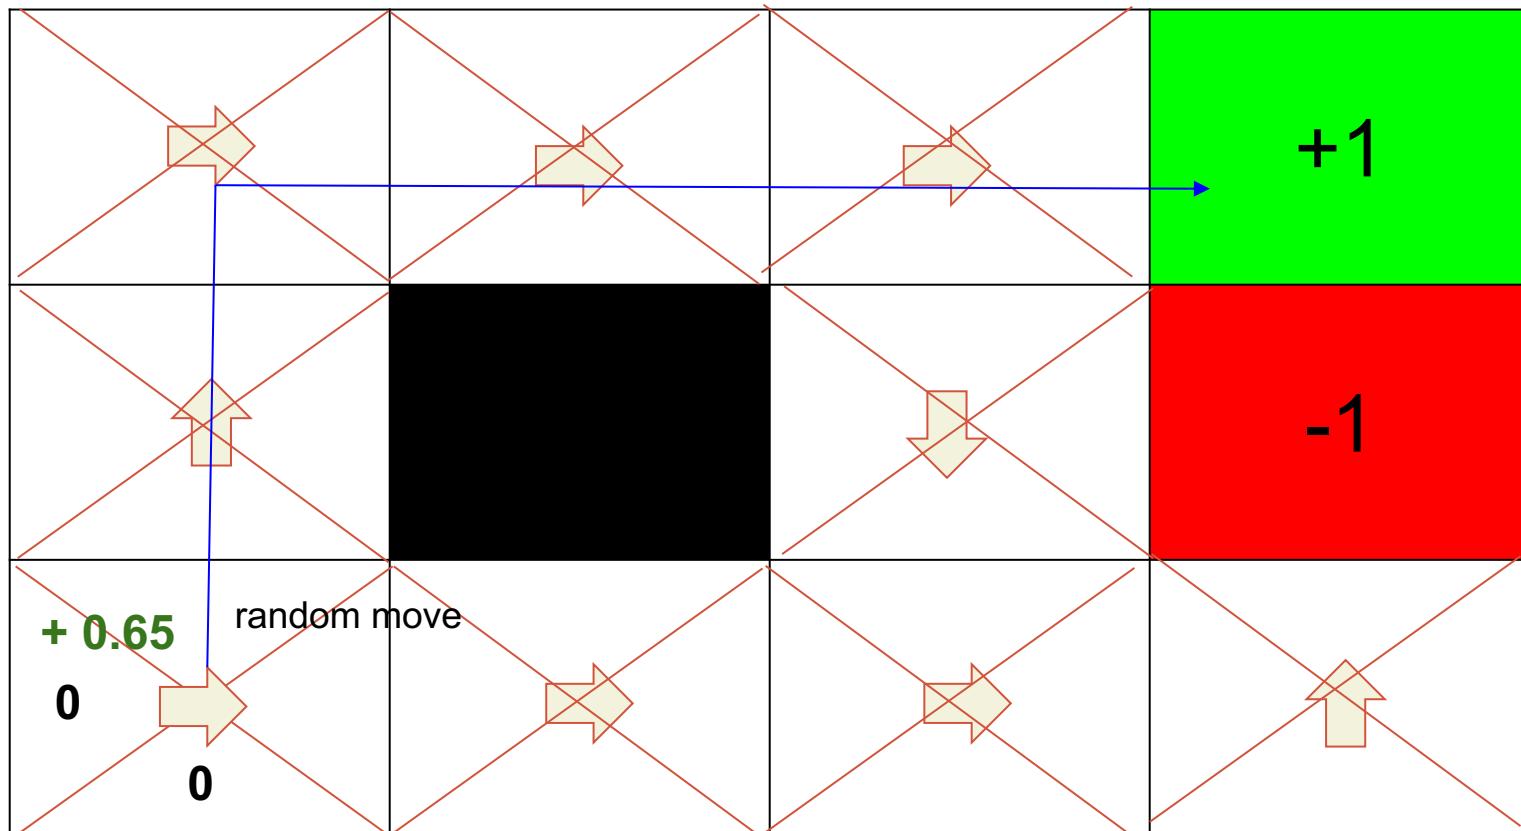
Running simulations until the end.
Estimate Q as the average of taking s,a



Exploration

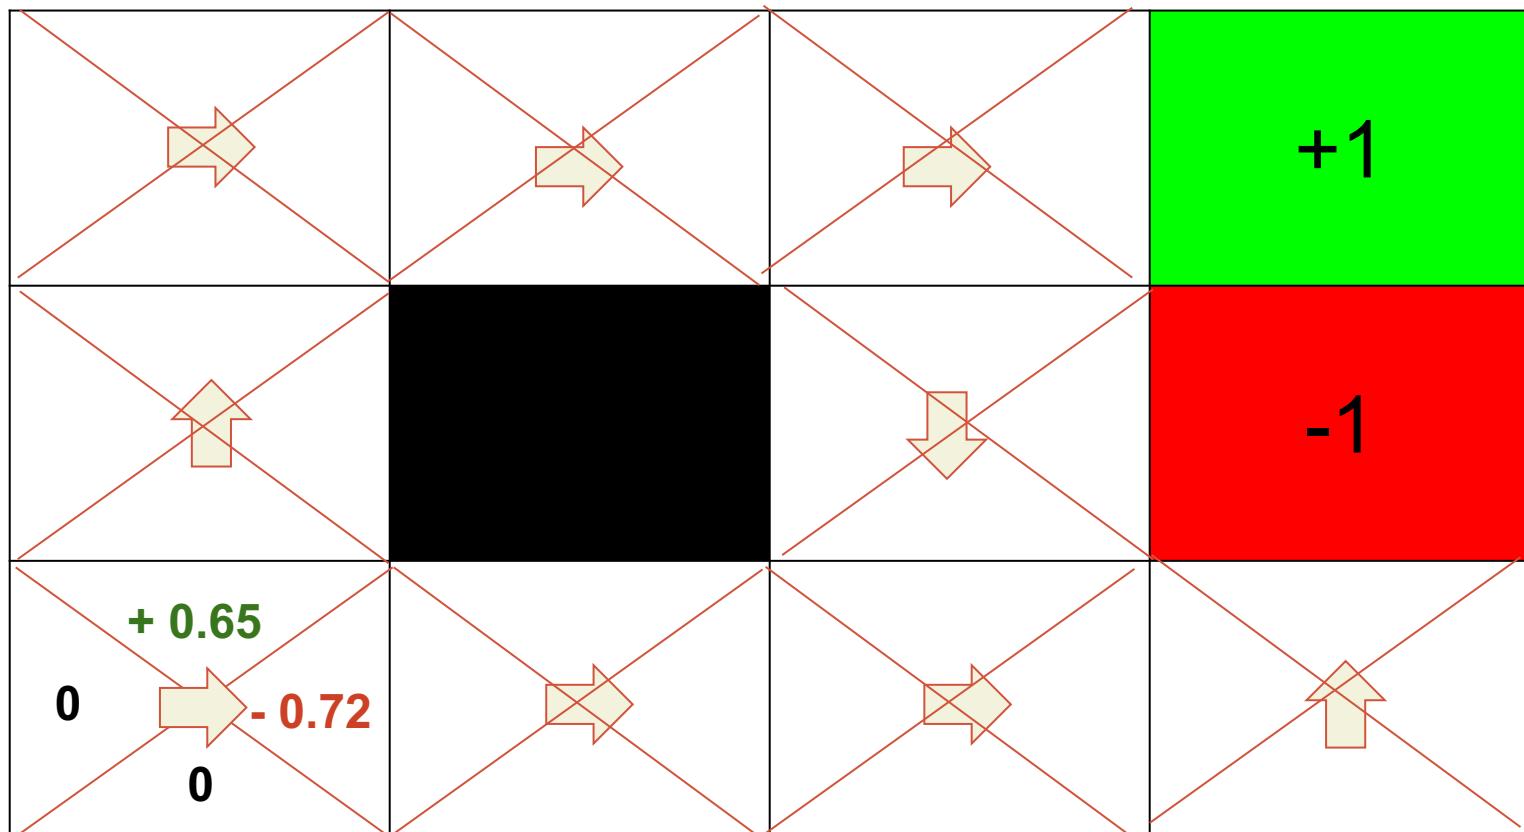
Randomly **explore** sometimes to get better solutions

Have epsilon probability to take a random move. Take current policy move with 1-epsilon probability.



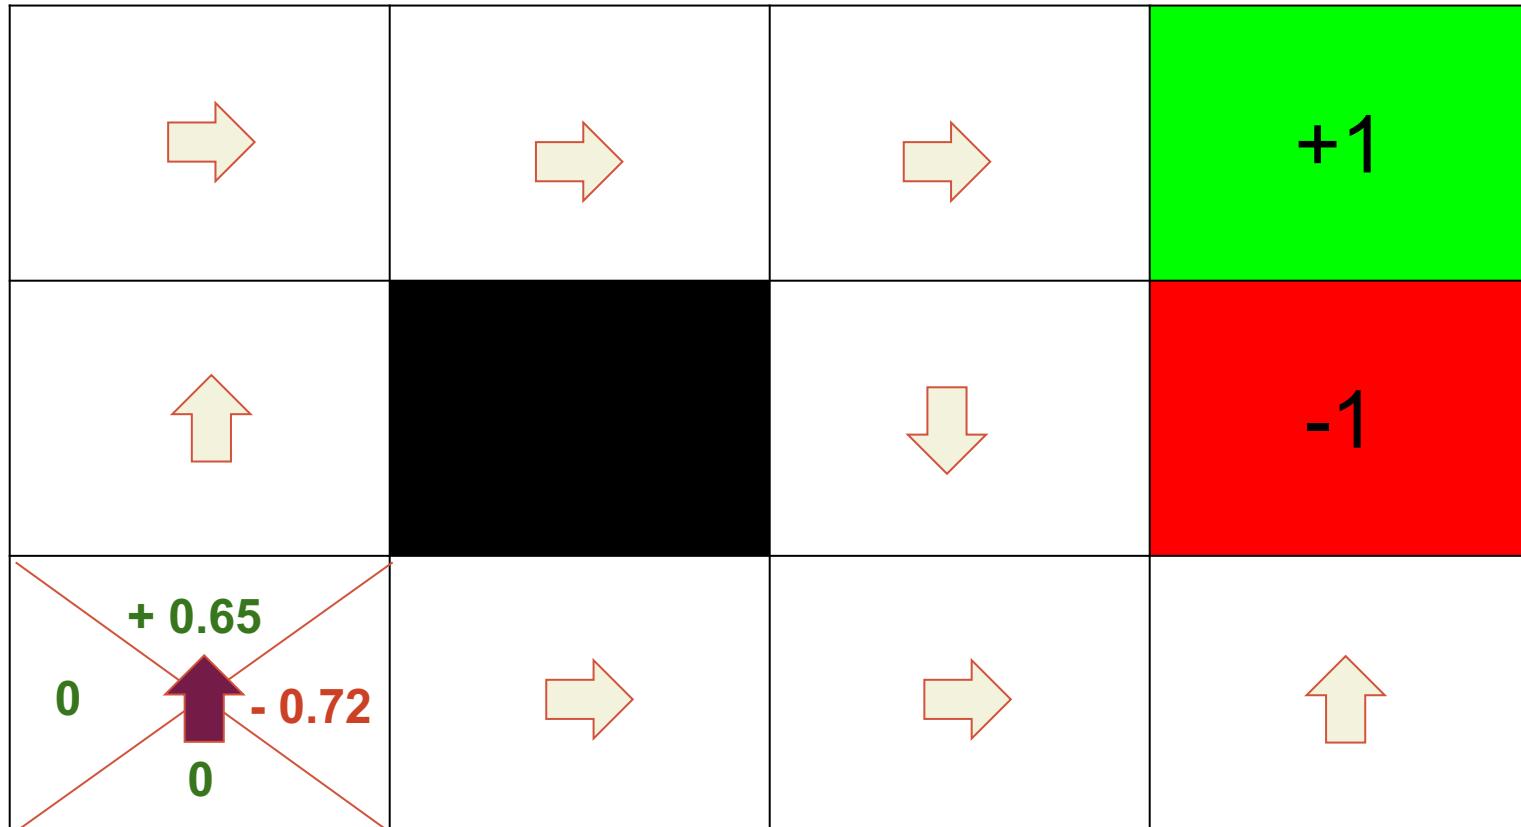
Monte-Carlo Estimator

Running many simulations until the end to estimate $Q(s,a)$



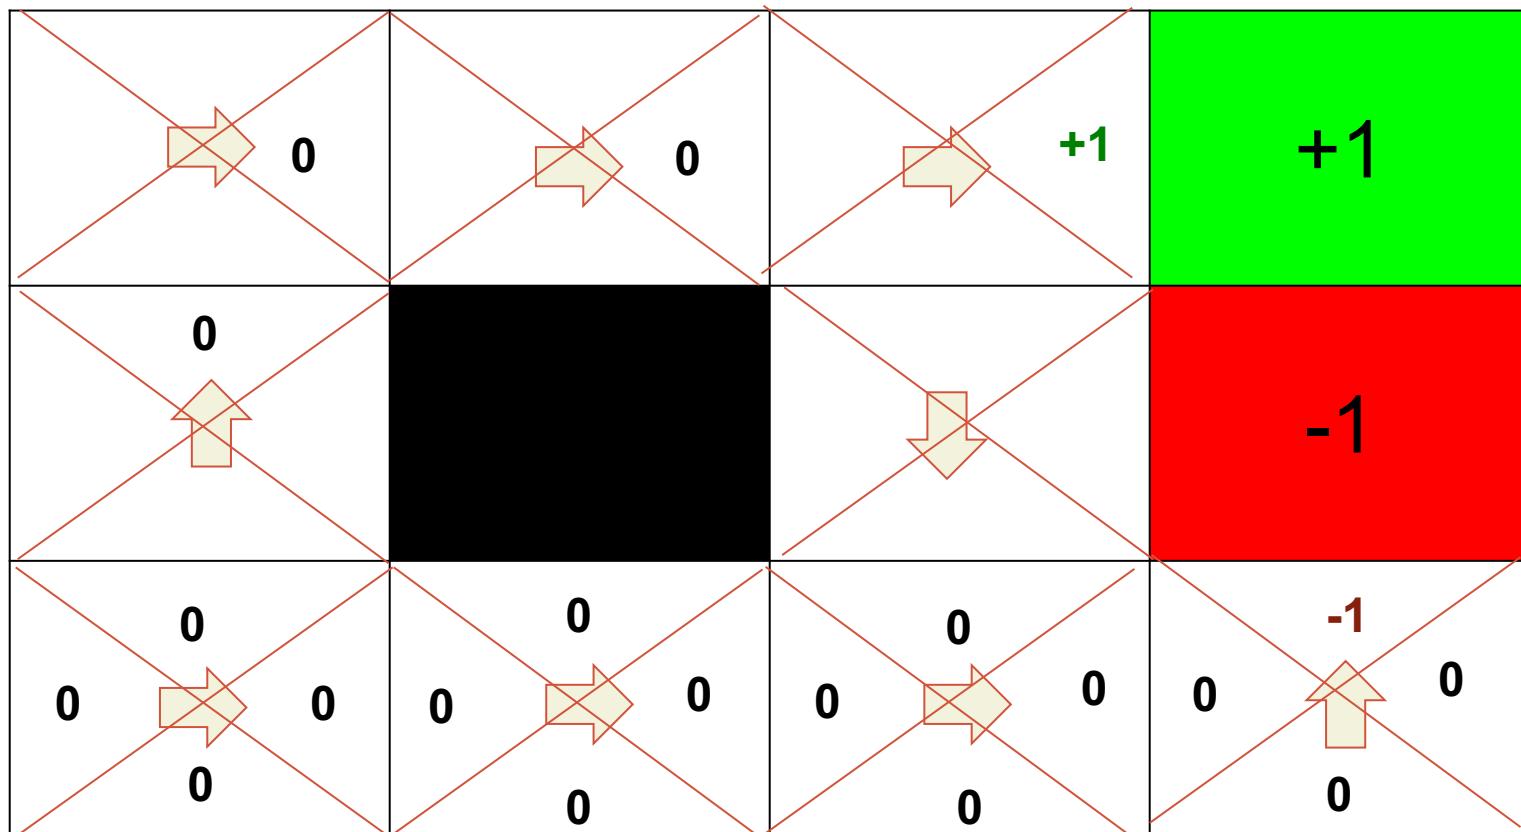
Policy Improvement

$$\pi'(s) = \operatorname{argmax}_a Q(s, a)$$



Bootstrap Estimator

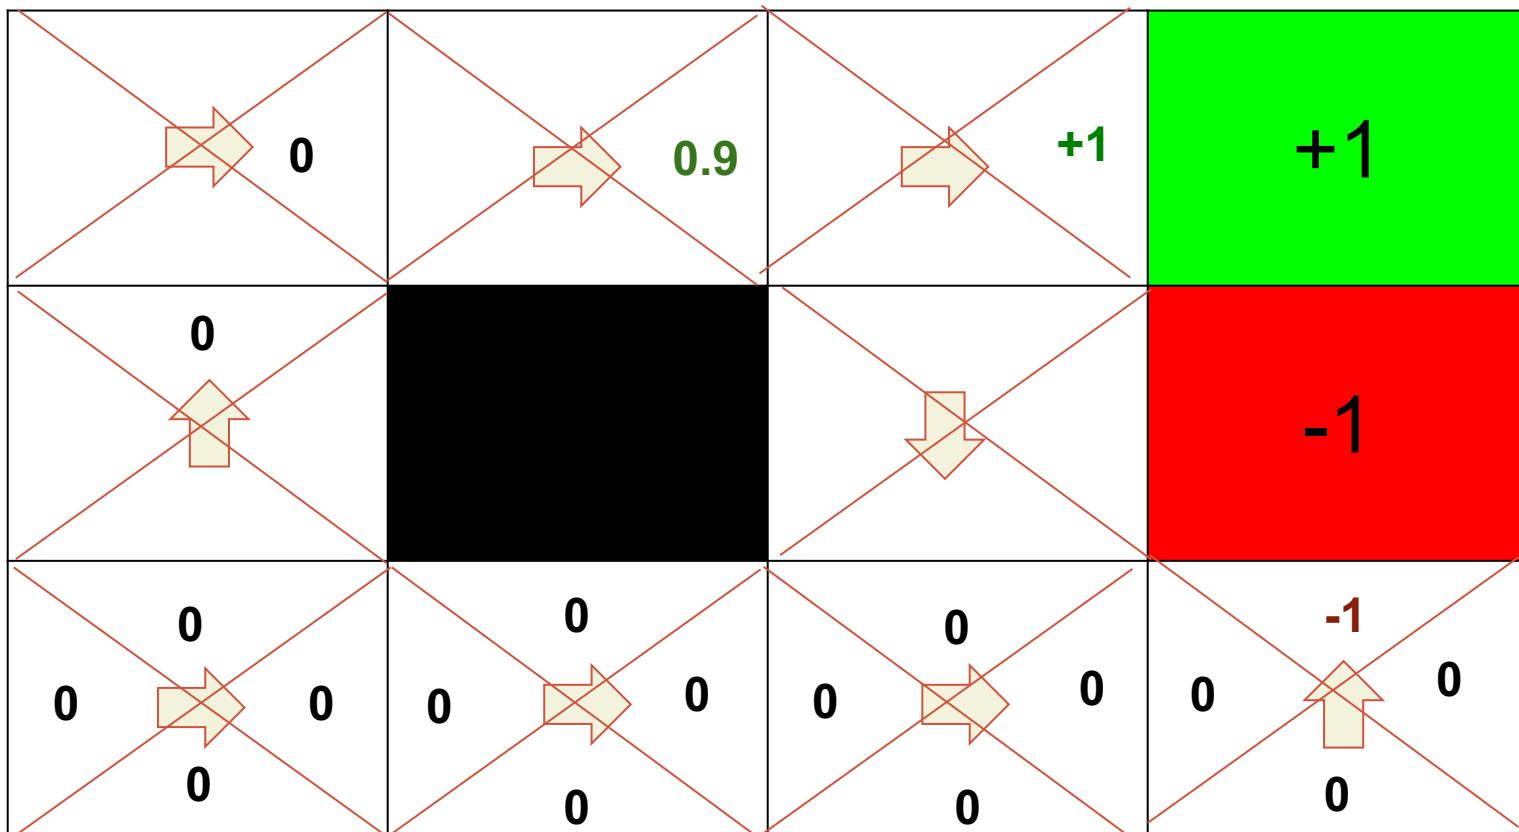
$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$



Bootstrap Estimator

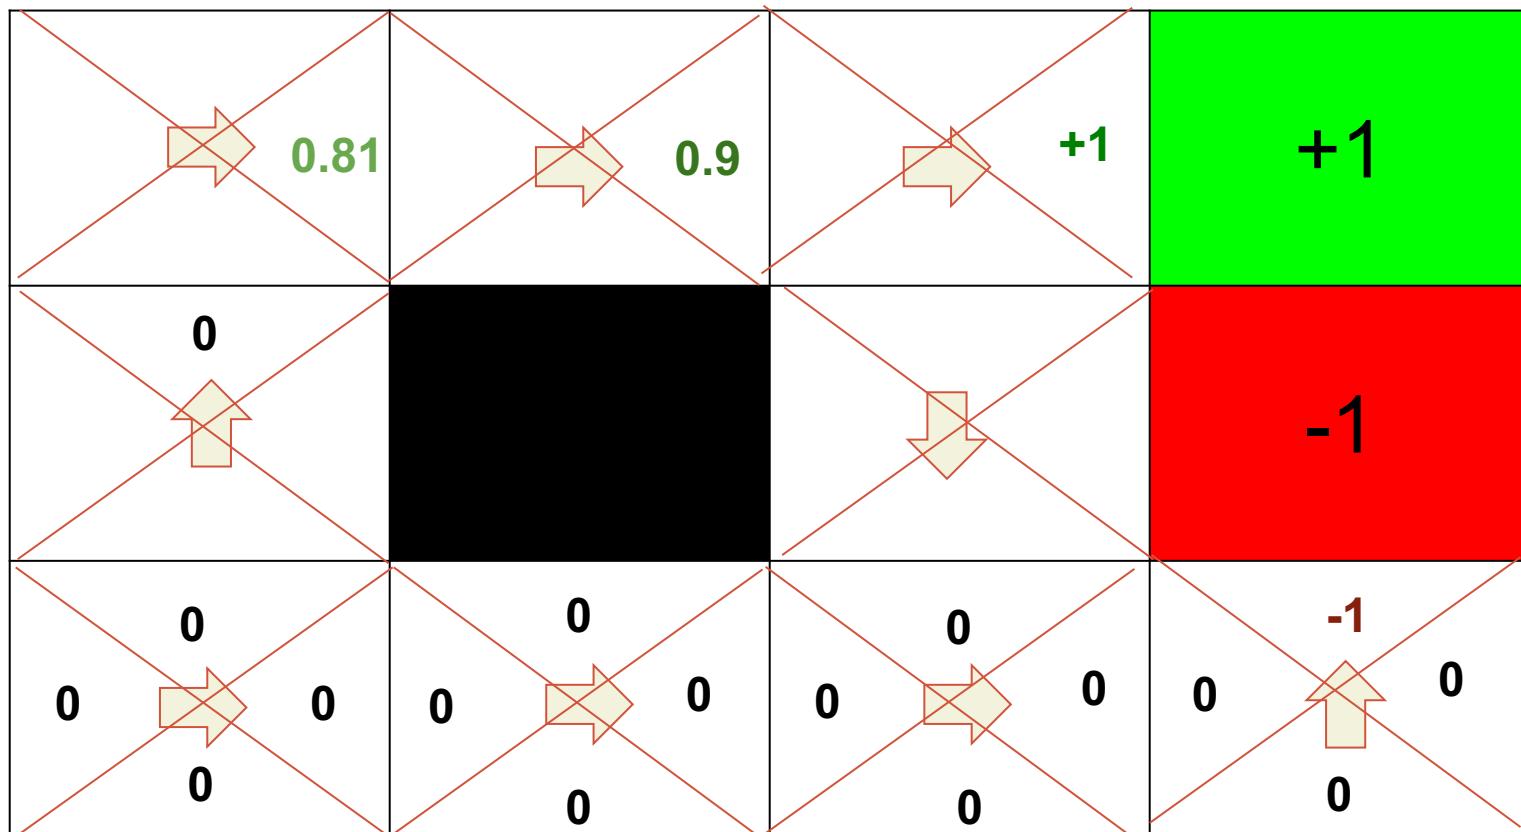
bellman equation

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$



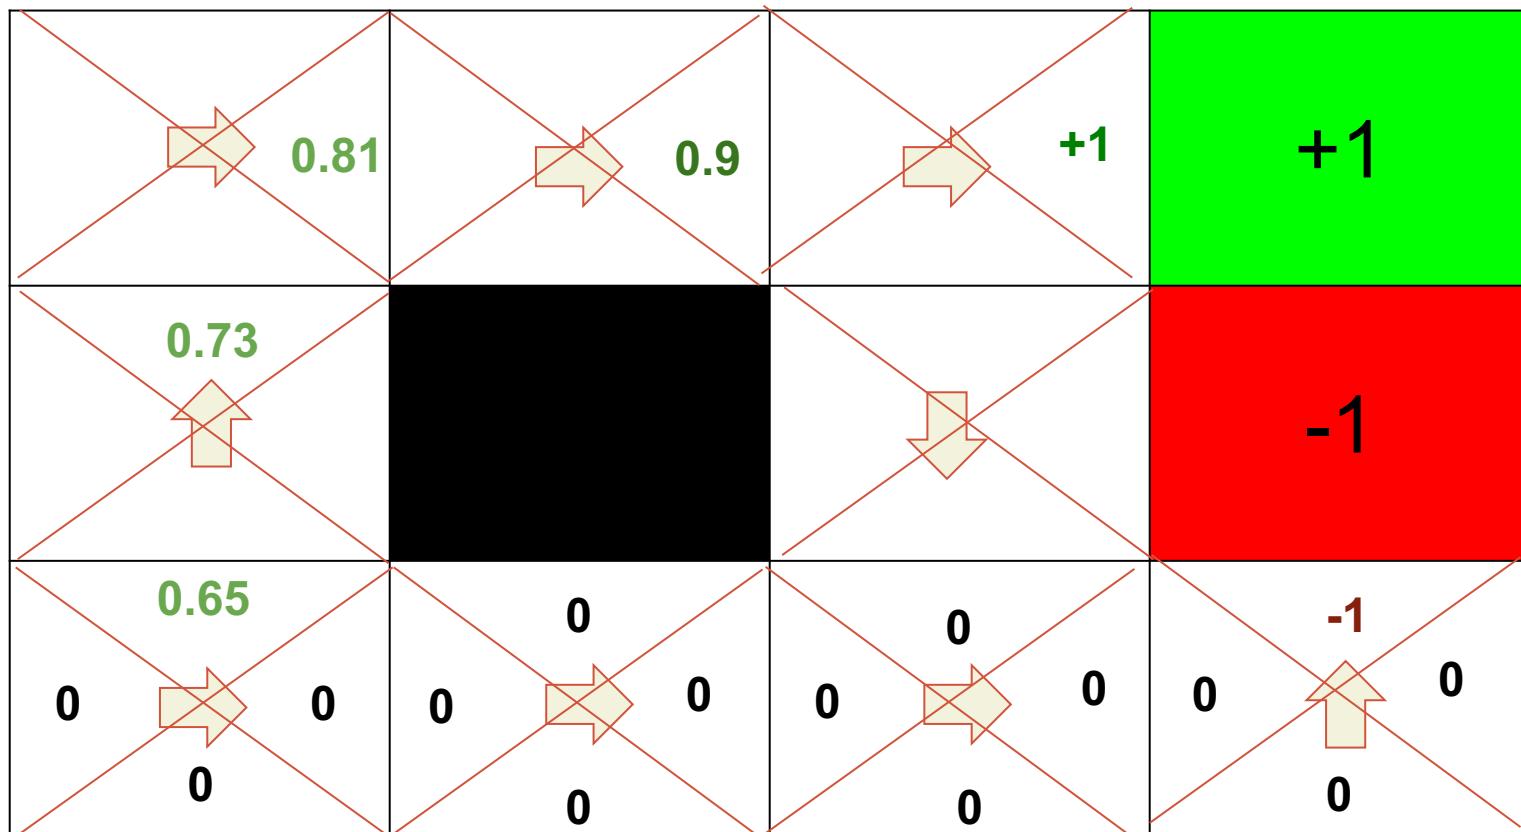
Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$

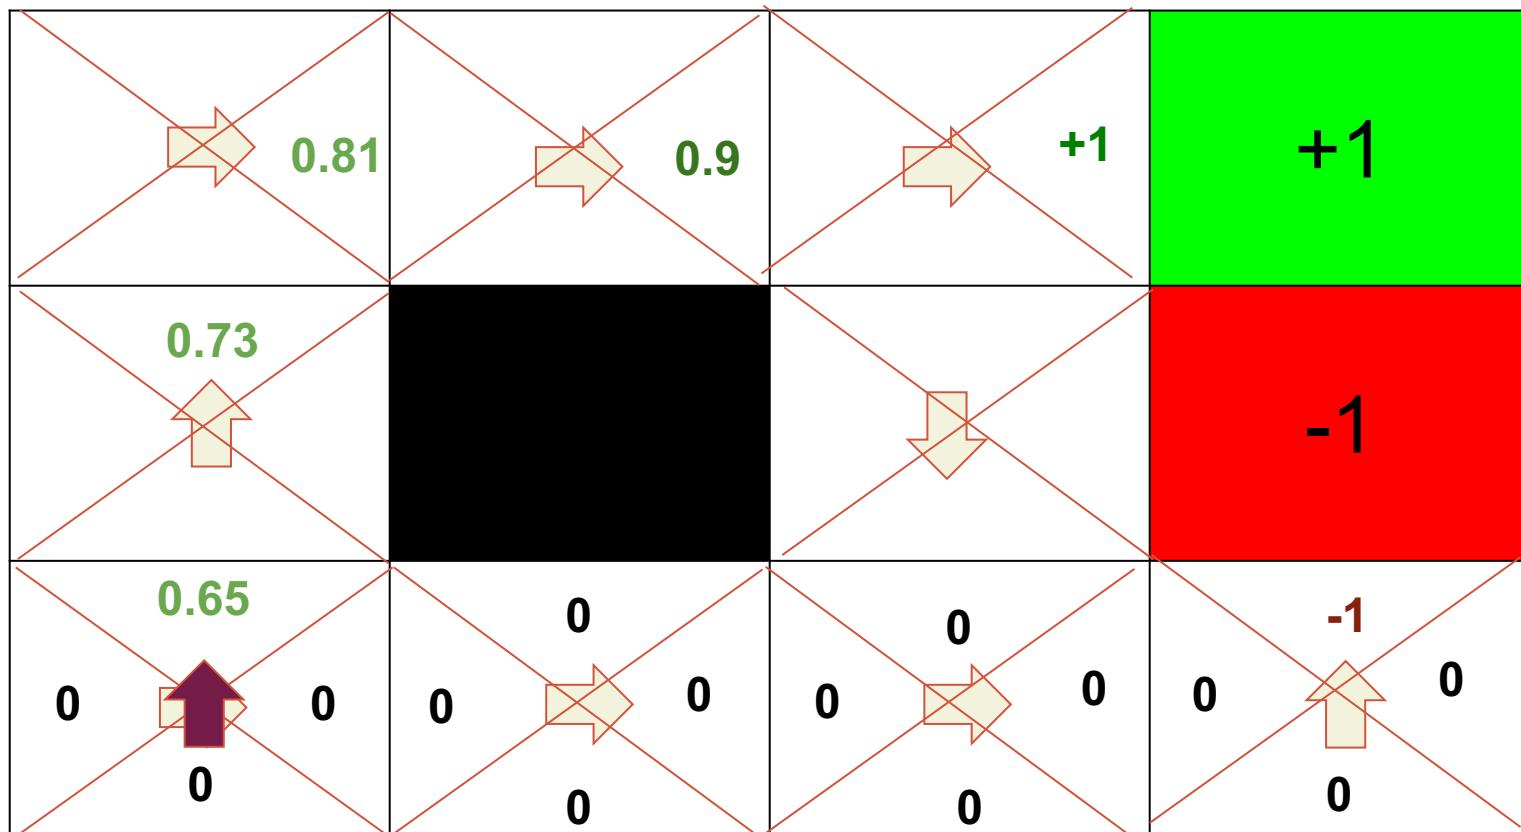


Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$



Policy Improvement



Bias and Variance in RL

What is bias of V_π estimation?

- Let $\hat{V}_\pi(s)$ be an estimate of $V_\pi(s)$
- $\hat{V}_\pi(s)$ is unbiased if:

$$\mathbb{E}_s \left[\hat{V}_\pi(s) - V_\pi(s) \right] = 0$$

What is variance of $\hat{V}_\pi(s)$ estimation?

$$\text{Var} \left[\hat{V}_\pi(s) \right] = \mathbb{E}_s \left[(\hat{V}_\pi(s) - \mathbb{E}_s [\hat{V}_\pi(s)])^2 \right]$$

- High if $\hat{V}_\pi(s)$ fluctuates a lot

Bias and Variance

- Monte-Carlo estimate has high variance and low bias.
- Bootstrap estimate has higher bias but lower variance.

Problems with RL

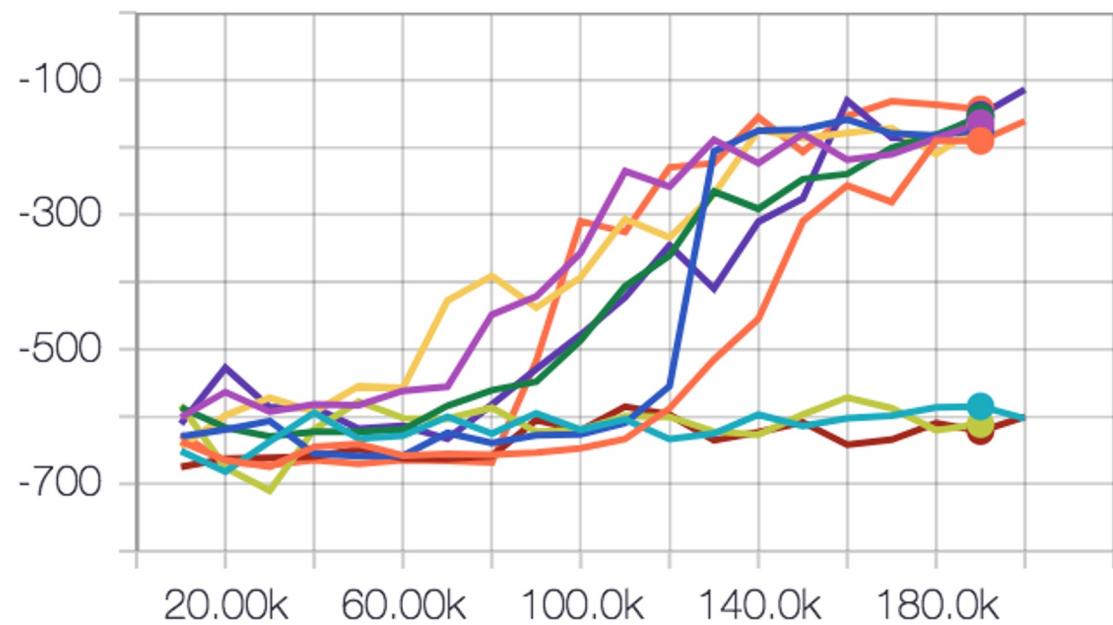
Randomness

Random initialization

Random exploration

Random environment

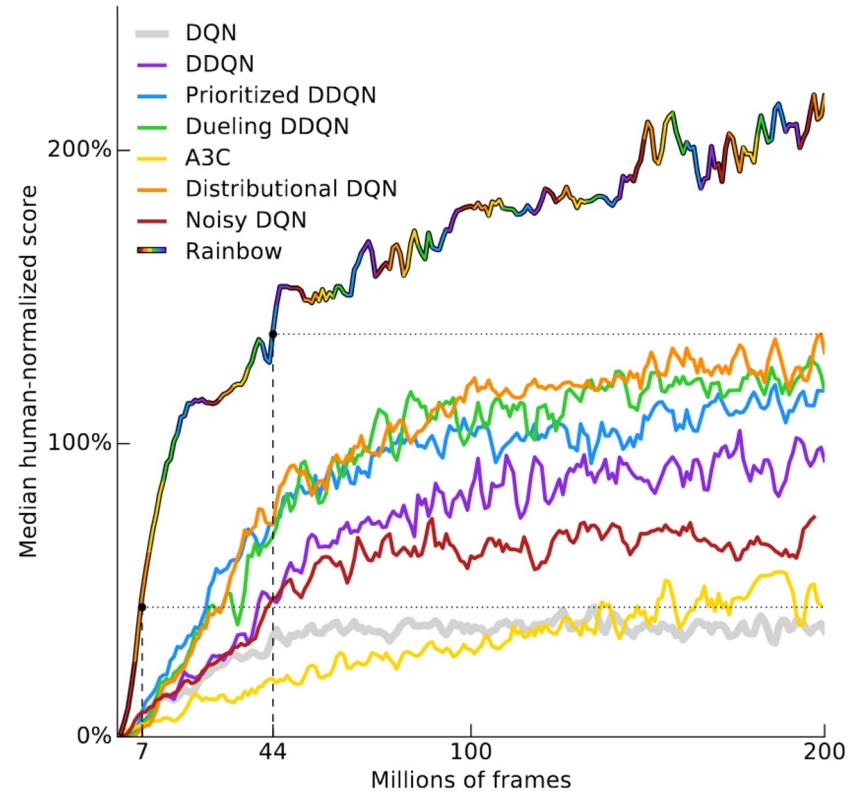
episode_reward/test



Problems with (current) RL

Data inefficient

- Many use case can be better solved with supervised learning (efficiency and accuracy)

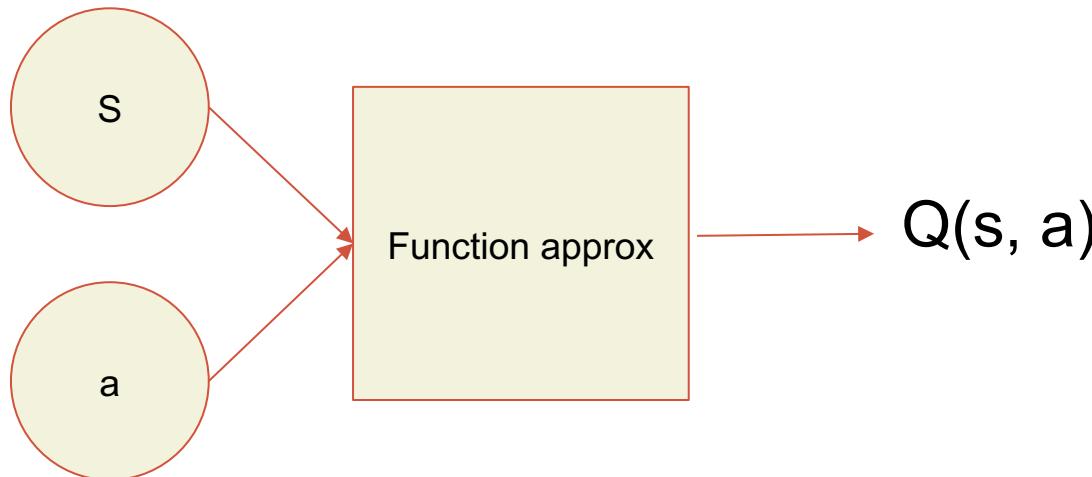
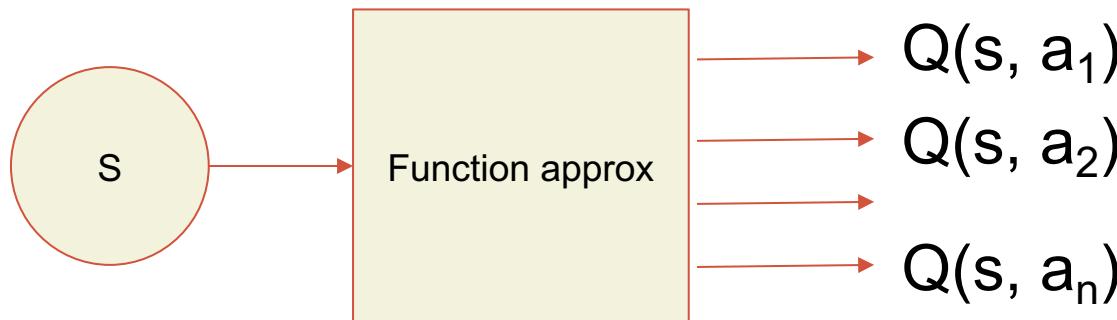


Function Approximator (FA)

- Tabular Q-value is impractical when the state-action space is large!
 - Need large memory
 - Impractical to fill up every cell
- Enter .. a function approximator

Function Approximator (FA)

- Instead of a table containing Q-value for every state and action, use a function that output Q-values.



Learning with FA

- With tabular Q-learning,
 - the act of learning = putting Q-value in the table
- With function approximator,
 - the act of learning = searching for the optimal parameters of the FA

Learning with FA

- How to adapt the parameters (weights) of the FA?
- Step 1: Define a loss function.
- Step 2: Optimise the weights to minimise the loss

Loss function

- What should be the loss function?
- Introducing Bellman's equations

$$V^\pi(s_t) = E_{\pi,P}[r_t + \gamma V^\pi(s_{t+1})]$$

$$Q^\pi(s_t) = E_{\pi,P}[r_t + \gamma Q^\pi(s_{t+1}, a_{t+1})]$$

- Bellman's optimality equations

$$Q^*(s_t) = E_P[r_t + \gamma \max_b Q^*(s_{t+1}, b)]$$

Loss function

- The Bellman's equation must hold for correct Q-value
- Rewrite the Bellman's optimality with our estimator (FA)

$$\hat{Q}(s_t) = E_P[r_t + \gamma \max_b \hat{Q}(s_{t+1}, b)]$$

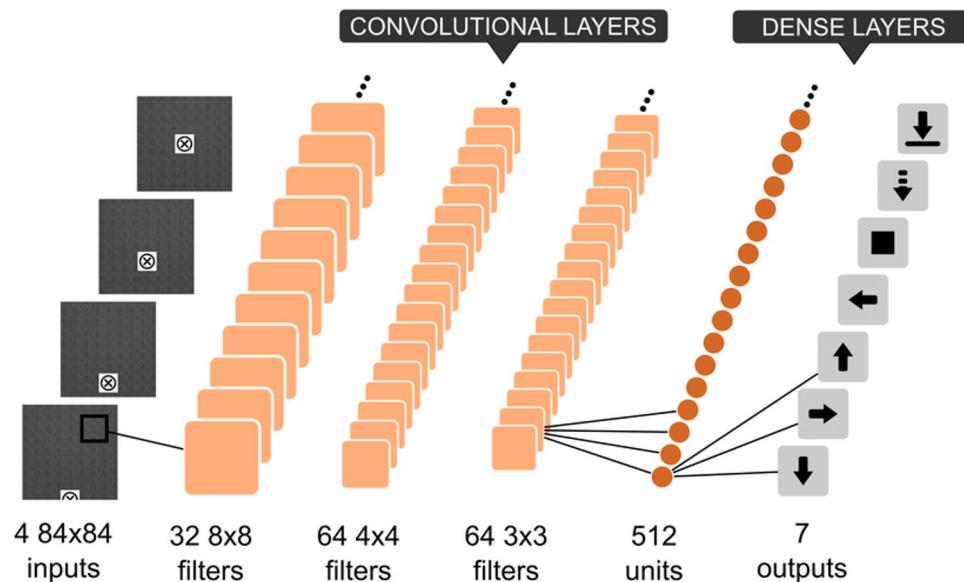
**The estimator is correct if
the left hand side = right hand side**

$$TD = r_t + \gamma \max_b \hat{Q}_\theta(s_{t+1}, b) - \hat{Q}_\theta(s_t)$$

``Temporal Difference error''

Temporal Difference Learning

- Use TD-error to guide learning
- Example
 - Deep Q-Networks (DQN)
 - Deep convolutional neural network as a function approximator
 - Optimise square TD-error



$$L(\theta) = (r_t + \gamma \max_b \hat{Q}_\theta(s_{t+1}, b) - \hat{Q}_\theta(s_t))^2$$

Policy-based methods

Policy gradient

Policy gradient

Q-Learning

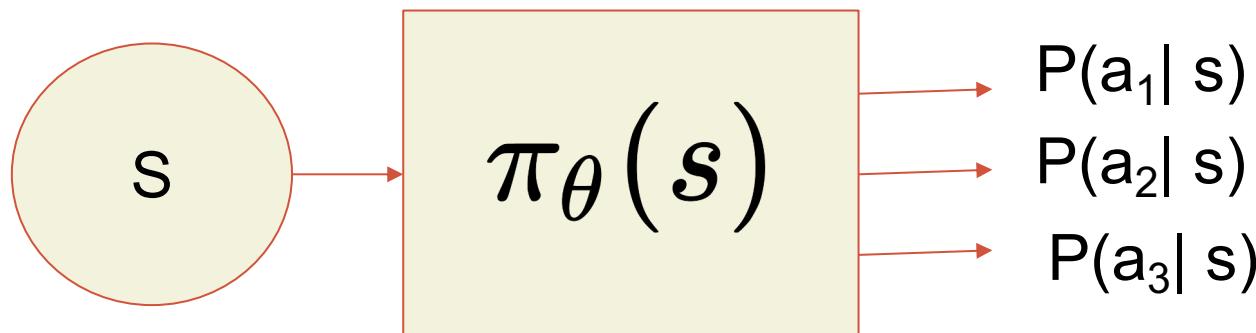
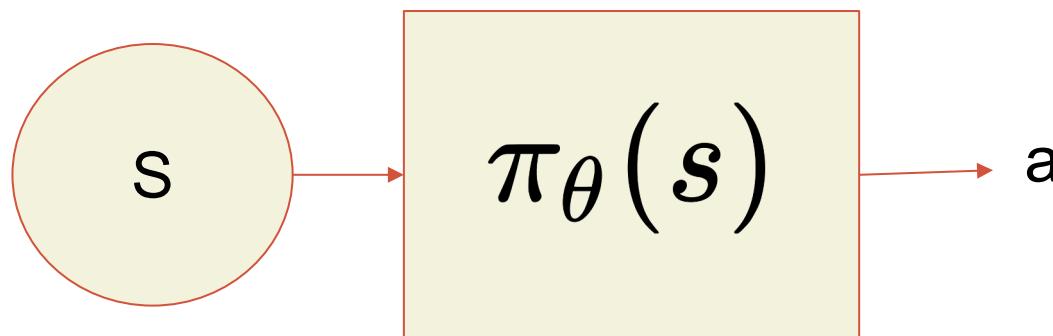
- policy is implicit
- if we already have Q, we have policy
- we just look at Q to get π

Policy gradient

- learns π directly explicitly
- use Q, V as a helper for learning π

Policy gradient

- Use Function Approximator to represent policy directly



Policy-based vs Value-based

Policy-based can learn continuous actions

Q-learning needs to $\text{argmax}_a Q(s,a)$

Policy-based can yield non-deterministic policy $P(a | s)$

Loss function for policy gradient

- Start-state objective

$$J(\theta) = E_{\pi(\theta)}[G|s_0] = V^\pi(s_0)$$

- Average-reward objective

$$J(\theta) = \sum_s d^\pi(s) \sum_a \pi(s, a) r(s, a)$$

* d is a stationary distribution of a Markov chain.

- One way to optimise these objectives is to use SGD.

Computing the gradient

Let's try to compute the gradient of the start-state objective

$$J(\theta) = E_{\pi_\theta}[G|s_0]$$

To evaluate this expectation, maybe we could try
a one-sample Monte-Carlo estimator:

$$J(\theta) \approx r_0 + \gamma r_1 + \gamma^2 r_3 + \dots$$

$$\nabla J(\theta) \approx \nabla_\theta [r_0 + \gamma r_1 + \gamma^2 r_3 + \dots] \quad \text{X}$$

Doesn't quite work? The evaluated value does not depend
on θ . Gradient can't be computed.

Computing the gradient

Maybe we can try change θ a little bit and find the difference?

$$J(\theta) \approx r_0 + \gamma r_1 + \gamma^2 r_3 + \dots$$

$$J(\theta + \delta\theta) \approx r'_0 + \gamma r'_1 + \gamma^2 r'_3 + \dots$$

$$\nabla J = \frac{J(\theta) - J(\theta + \delta)}{\delta}$$

Could work? But...

Looks very expensive and noisy to compute!

Maybe there is a better way?

Policy gradient

Let's start from the average-reward objective

$$J(\theta) = \sum_s d^\pi(s) \sum_a \pi_\theta(s, a) r(s, a)$$

For simplicity let's assume $d(s)$ does not depend on θ

$$J(\theta) = \sum_s d(s) \sum_a \pi_\theta(s, a) r(s, a)$$

$$\nabla_\theta J(\theta) = \sum_s d(s) \sum_a \nabla_\theta \pi_\theta(s, a) r(s, a)$$

Almost there...

Policy gradient

REINFORCE trick!

$$\nabla_{\theta} J(\theta) = \sum_s d(s) \sum_a \nabla_{\theta} \pi_{\theta}(s, a) r(s, a)$$

$$\nabla_{\theta} J(\theta) = \sum_s d(s) \sum_a \pi_{\theta} \nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)$$

We get this by sampling a playthrough using current policy (on-policy)

$$\nabla_{\theta} J(\theta) = E_{\pi} [\nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)]$$

$$\nabla_{\theta} J(\theta) \approx \nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)$$

Policy gradient theorem

There is a theorem...called policy gradient theorem
say that we can replace $r(s, a)$ with $Q(s, a)$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)]$$

What is the gradient doing?

Goal maximize rewards

Push π towards directions of higher $Q(s,a)$

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)]$$

$$Q(1,1) = 5$$

$$Q(1,2) = 2$$

▼ $\pi(1,2)$ will have a higher weight. Policy gets push towards action 2

Notes on Policy gradient

Also known as REINFORCE or likelihood ratio.

Used by other ML fields when original loss is not differentiable (-Q in this case). Push network to produce lower loss.

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [Q^{\pi_{\theta}}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)]$$

Example of non-differentiable functions

argmax (not maxpool)

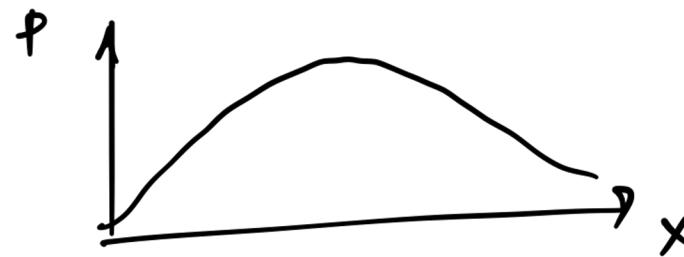
sampling

Encourage Exploration

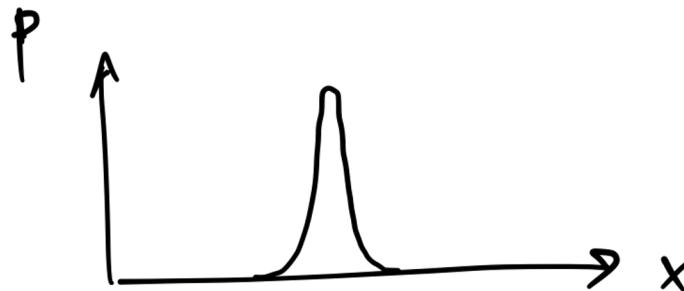
- policy $\pi_\theta(a|s)$ could be too confident early
- Like, $\pi_\theta(a = a|s) = 1$
- This could lead to insufficient exploration
- Encourage exploration by “entropy term” $H(\pi_\theta)$
- We want to punish too low entropy
- New gradient rule: $\nabla_\theta J(\theta) \rightarrow \nabla_\theta J(\theta) + \nabla_\theta H(\pi_\theta)$

Entropy (H)

high entropy



low entropy



On policy and off policy algorithms

Q Learning: $Q^*(s_t, a_t) = r_{t+1} + \max_a Q^*(s_{t+1}, a)$

- you need a_t , s_t , r_{t+1} , s_{t+1} to satisfy the above equation
- you can get (a, s, r, s') from any policy
- Q learning is *off-policy*

Policy gradient: $\nabla_\theta J(\theta) = \mathbb{E} [Q_\pi(s, a) \nabla_\theta \log \pi_\theta(a|s)]$

- you need s , a and $Q_\pi(s, a)$
- Q_π needs to be from the current policy
- s , a needs to come from current policy
- policy gradient is *on-policy*

Notes on on vs off policy

Off-policy can learn from any policy

- Can use old experience

- Can use expert experience

- Sample efficient

On-policy can only learn from current policy

- Slow

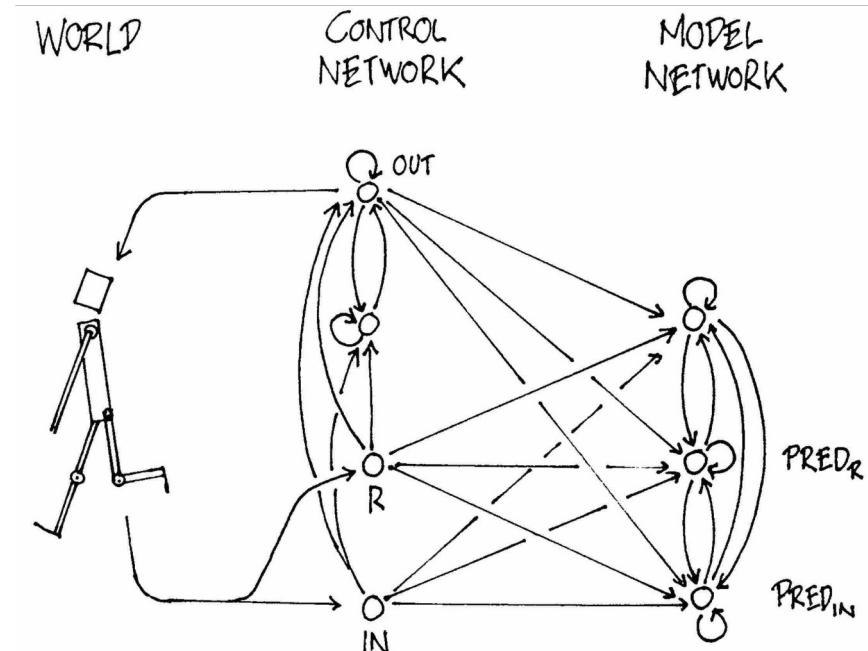
Model-based RL

RL's Model

RL's Model vs ML's Model

Think mental model that models the environment

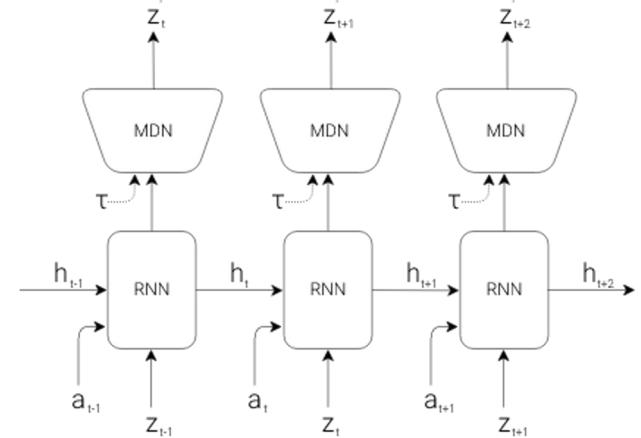
Basically a model can guess the future or we have access to the environment so we can take alternative routes



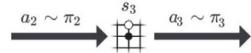
RL's Model

Guess the future

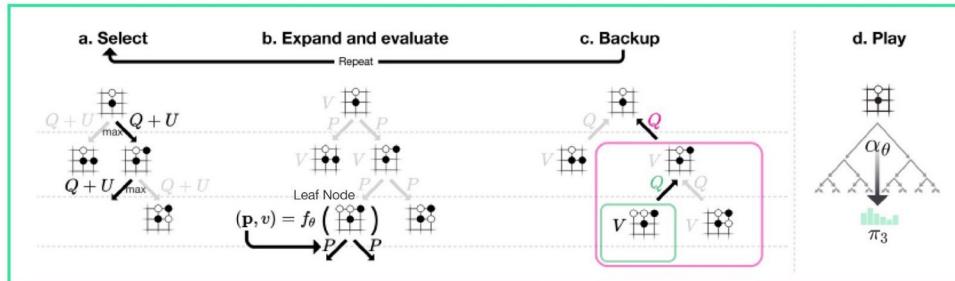
Or access to the environment
so we can take alternative
routes



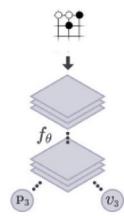
With a board position s_3



Use MCTS
to compute π_3



Use f to compute
 p and v



$$u(s, a) = c_{\text{pert}} P(s, a) \frac{\sqrt{\sum_b N(s, b)}}{1 + N(s, a)}$$

$$(P(s, \cdot), v) = f_\theta(s)$$

$$N(s, a) = \sum_{i=1}^n I(s, a, i)$$

$$W(s, a) = W(s, a) + v$$

$$Q(s, a) = \frac{W(s, a)}{N(s, a)}$$

$$\pi(a | s) = N(s, a)^{1/\tau} / \sum_b N(s, b)^{1/\tau}$$

<https://worldmodels.github.io/>

https://medium.com/@jonathan_hui/monte-carlo-tree-search-mcts-in-alphago-zero-8a403588276a

Model-based RL

- In model-based RL, we first build the model of the environment
- Then use that model to directly search for the answer.
- The problem is ... inaccurate model can give us bad policies...
- It is believed that if we can treat the uncertainty in the model correctly...model-based RL is the most efficient method!
- However, measuring uncertainty in the model is also very difficult.

Things to consider

When do we need RL?

- Your action affects the observation. Action has consequences (RL vs Bandit problem)
 - $x_1, a_1 \xrightarrow{\text{orange arrow}} x_2$
- The target behavior is difficult to be directly hard-coded.
 - How to move a snake robot?
- Collection of the data of a target behavior is difficult.

AlphaGo and why it works so well

Properties of the game Go

Deterministic

Fully observable

Rules are known (Model completely known)

Static

Fits MDP (Markov property)

Credit assignment problem

- An action can have consequences further away in time
- Some movements might not have any effect on the outcome



$$r_t = 0$$



10 time steps later



$$r_{t+10} = 1$$



Sparse reward problem

- Another problem is when rewards are sparse.
- Since model-free RL is just learning the correlations of trajectories and rewards... when there is no reward, RL cannot learn.
- Can we make it better?
 - Curiosity + intrinsic motivation?
 - Curriculum learning?
 - Hierarchical RL?

Designing the reward signal

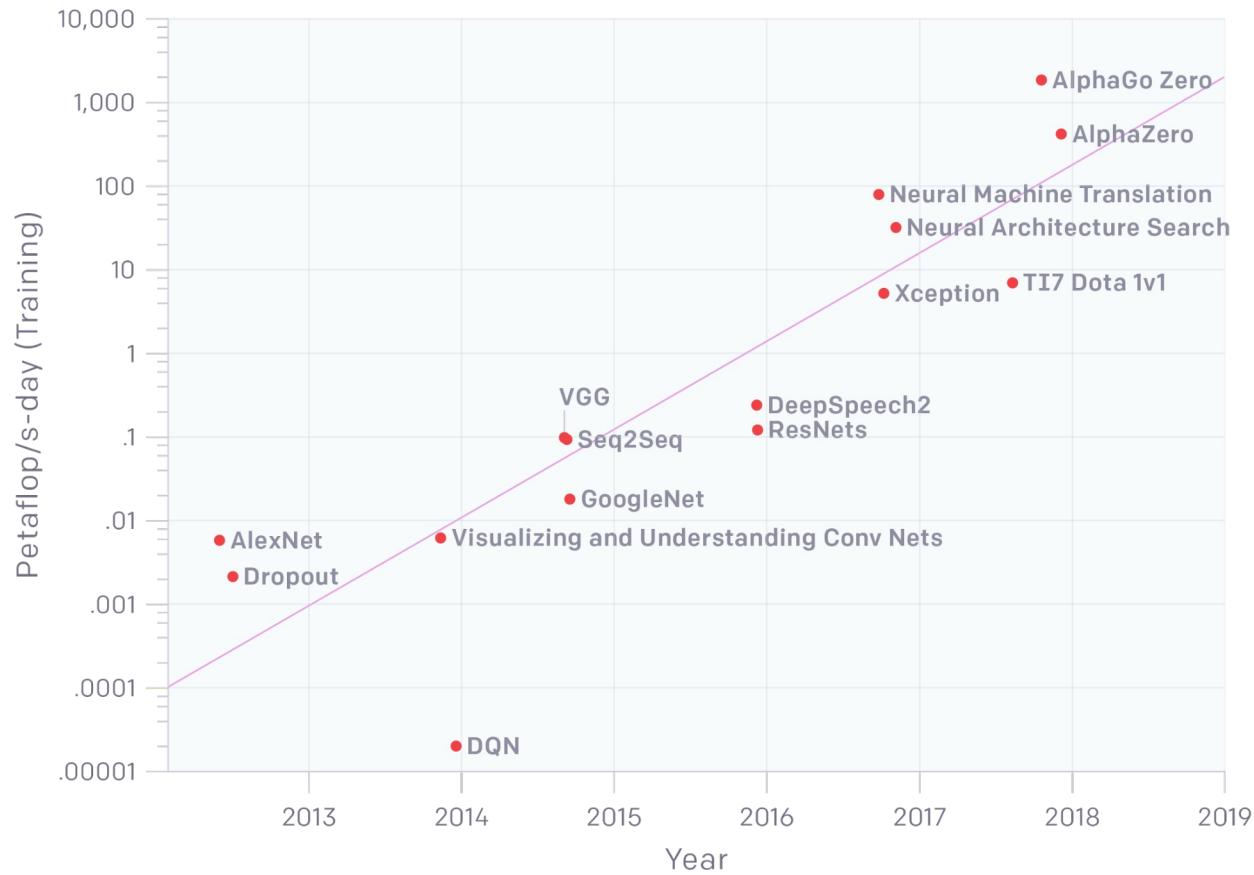
- Reward design can be quite challenging..
- Naive reward design can lead to unexpected (cheating) behaviours!
- Example:



<https://www.youtube.com/watch?v=tIOIHko8ySg>

Problems with (current) RL

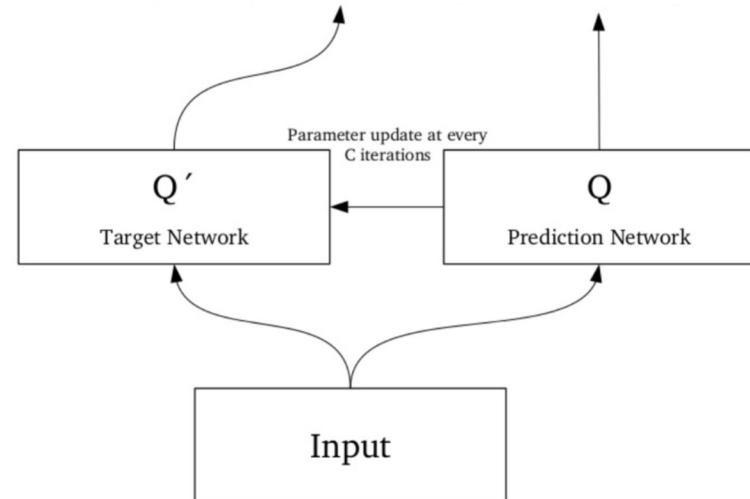
AlexNet to AlphaGo Zero: A 300,000x Increase in Compute



Bias and Variance

- Bias and variance are really important in RL
- We want to reduce both of them as much as possible
 - DQN uses experience replay to reduce bias
 - DQN uses target network to reduce variance
 - Actor-Critic method uses baseline to reduce variance
 - Actor-Critic method uses parallel worker to reduce bias
 - etc.

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$



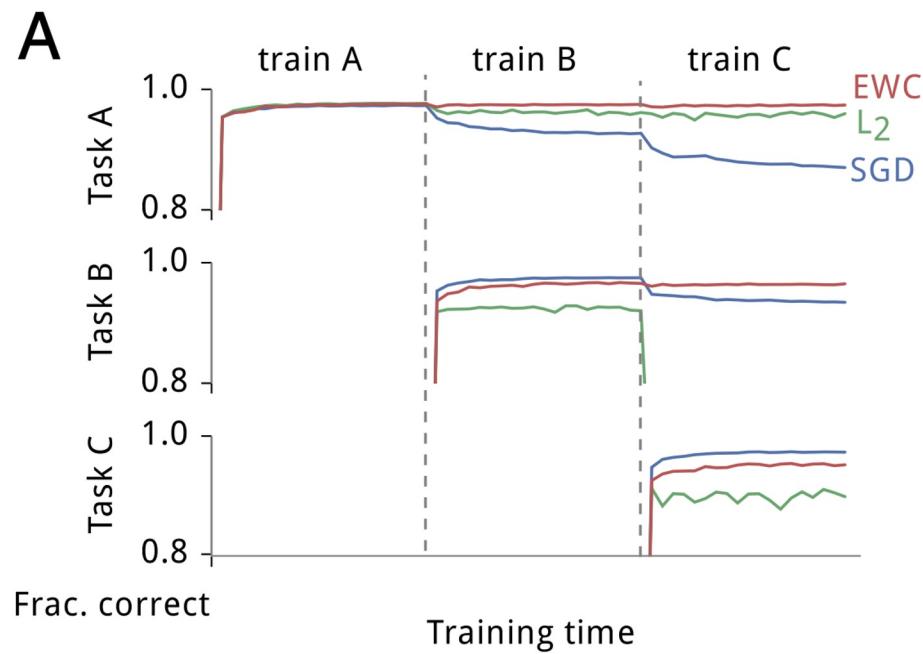
<https://www.slideshare.net/MuhammedKocaba/human-level-control-through-deep-reinforcement-learning-presentation>

Exploration and Exploitation

- Many RL results assume that all states are visited infinitely often.
 - Also, many RL algorithms are reduced into just an optimization problem.
 - Therefore, nicely spread/informative data can help a lot!
-
- DQN uses epsilon-greedy for exploration
 - Policy gradient uses entropy regularizer to encourage exploration

Optimisation problem

- Initialization problems
- Is SGD the best we can do?
- Catastrophic forgetting?

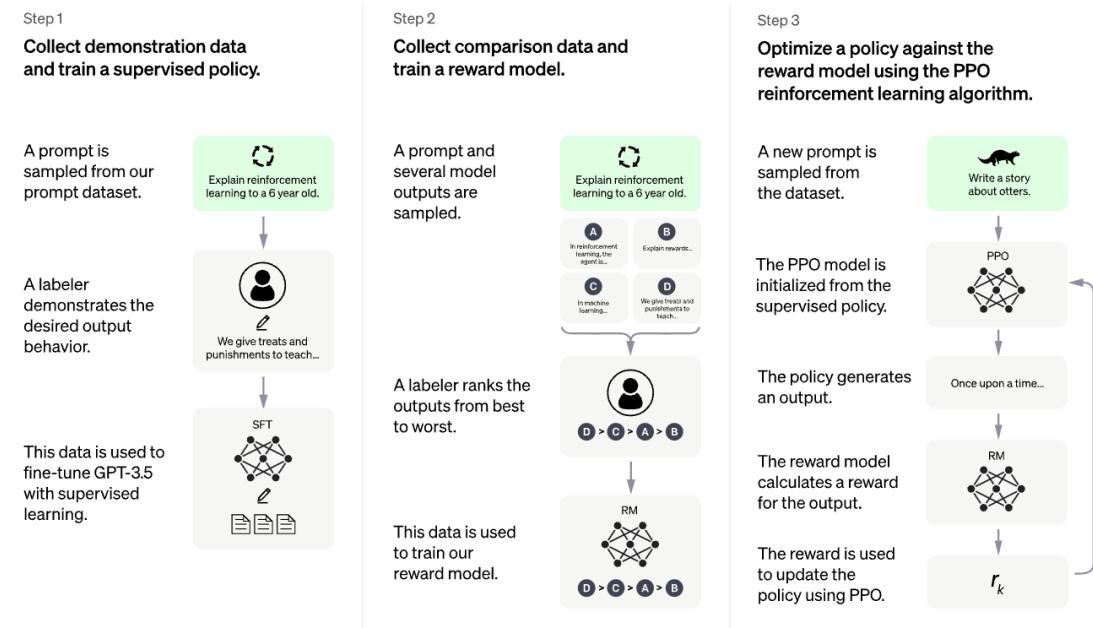


Current trends & open problems

- Intrinsic motivation, reward-bonus
- Imitation learning
- Multi-agent system and self-play
- Curriculum learning
- Model-based RL
- Robot learning + sim-to-real transfer learning
- etc...

RL and ChatGPT

- ChatGPT was trained using RLHF (Reinforcement Learning with Human Feedback). It uses PPO, a state-of-the-art policy gradient method.
- Recent works start to question whether we really need to use RL (<https://arxiv.org/abs/2305.18290> <https://arxiv.org/abs/2403.07691>)



Reinforcement learning

Elements of RL

- Environment, Agent, State, and MDP

Estimating Q

- Monte-Carlo, Bootstrap

- Deep learning as a function approximator

Policy learning

- Q-learning

- TD learning

- Policy gradient

Concepts

- Exploration vs Exploitation

Further learning

- Chula RL course
 - <https://www.youtube.com/playlist?list=PLcBOyD1N1T-PyNUNA77ITYNCAeAMGxV5I>
- Deepmind RL course 2021
 - <https://www.deepmind.com/learning-resources/reinforcement-learning-lecture-series-2021>